

You hit on a fundamental truth of how AI works: LLMs are inherently stateless. They completely lose their memory the second your session ends.

Here is how you handle persistence, lead agent configuration, and those YAML blocks in Claude Code.

1. Can you configure the Lead Agent?

Yes, but it uses a different format than the subagents.

Subagents use those specific Markdown files in the `.claude/agents/` folder. The **lead agent** (your main session) is configured via a JSON file, typically located at `~/.claude.json` (for global settings) or `.claude.json` in your project folder. This is where your default model, theme, and API keys are stored permanently so you don't have to reset them every time.

2. How Persistent Memory Works

You are exactly right on the architecture — because the agent loses its memory, you *have* to give it an external file to read from every time it wakes up.

In Claude Code, you don't need a complex database. It handles persistent memory using a single file called **CLAUDE.md**.

Think of CLAUDE.md exactly like that "Wiki MD file" you mentioned. You drop it in the root directory of your project. Every time you start a new Claude Code session, the lead agent silently reads CLAUDE.md before it even says hello.

- **What goes in it?** Build commands, architectural rules, testing conventions, and project context.
- **What about JSONL caches?** If you were building your own custom on-device AI from scratch, structuring all your data into a localized JSONL cache is a brilliant way to give the agent super-fast reference capabilities. But out-of-the-box, Claude Code handles this by just running fast file searches (via `grep` and `glob` tools) over your raw text and code files whenever it needs to recall something not in CLAUDE.md.

3. What is the YAML Block?

The "YAML block" (often called YAML frontmatter) is just a standard way to attach machine-readable settings to a human-readable text document.

It is always laid out the exact same way: it must be at the **very top** of the file, and it must be sandwiched between two lines of three dashes (`---`).

Here is the exact layout:

```
YAML
---
name: my-agent
description: What this agent does
```

```
model: haiku
```

```
tools: Read, Bash
```

```
--
```

```
Everything below the second set of dashes is the system prompt (the instructions the agent actually reads).
```

The system orchestrator reads the top part (the YAML) to know *how* to run the agent, and the AI model reads the bottom part (the Markdown text) to know *what* to do.