

I hear you. First of all, the reason most AI agents completely shit the bed in Termux is because they assume standard Linux file paths. They look for `/usr/bin` or `/bin/bash`, completely failing to realize that Termux runs in a sandboxed, synthetic root environment (`/data/data/com.termux/files/usr/`). If you don't explicitly enforce Termux paths in the system prompt, they will break your environment trying to run commands that don't exist.

As for the pricing—you are 100% right. Running an autonomous agent like Claude Code directly on the Anthropic API is a fast track to draining your bank account because you pay per token for every single file it reads and writes.

Here is exactly how you bypass the API fees, hook up Unsloth, and get Aider ("eighter") or VS Code running on a flat budget or totally free.

1. The 100% Free Local Route: Unsloth + Aider + Hermes

Unsloth is incredible because it isn't just for fine-tuning anymore; **Unsloth Studio acts as a local API server**. You can load up Hermes (or whatever model you want) in Unsloth, and it will serve it up just like OpenAI or Anthropic, but running locally on your hardware.

How to connect it:

1. **Run Unsloth Studio** and load your model (like Hermes or Open Wiki LLM).
2. **Generate a Local API Key:** In the Unsloth UI, go to Settings → API Keys and generate one (it starts with `sk-unsloth-`).
3. **Get the Local URL:** Unsloth will expose an endpoint, usually something like `http://localhost:8888/v1`.
4. **Point Aider at Unsloth:** When you launch Aider in your Termux shell, you just tell it to use your local Unsloth server instead of the cloud. You run it with standard OpenAI-compatible flags:
`aider --openai-api-base http://localhost:8888/v1 --openai-api-key sk-unsloth-... --model <your-hermes-model-name>`

Now Aider is doing the heavy lifting in your terminal, and Unsloth is handling the local inference. Total API cost: **\$0**.

2. The Flat-Rate "Pro Account" Route (VS Code / Cursor)

When you mentioned getting a "Pro account" in VS Code that is way cheaper, you are likely thinking of a few specific setups that bundle API access into a flat monthly fee (usually around \$20/month) so you never pay per token.

Here are the best ways to get that flat-rate access without paying raw Claude API fees:

- **Cursor (The VS Code Fork):** This is probably what you were thinking of. Cursor is built on VS Code, and its \$20/month Pro tier gives you flat-rate access to top-tier models (including Claude Opus/Sonnet and GPT-5). It has a built-in agent that reads your codebase, and you never have to plug in a raw Anthropic API key.
- **Cline or Continue.dev (VS Code Extensions):** If you don't want to use Cursor and want to stay in standard VS Code (or an IDE running via Termux), install an extension like **Cline**.

Cline allows you to "Bring Your Own Key" (BYOK). You can point Cline directly at your Unsloth local server, giving you an elite VS Code agent completely for free, powered by your local Hermes model.

- **GitHub Copilot Pro / Codex:** Copilot now has robust agent modes in VS Code for \$10 to \$20 a month, and it includes access to Claude and GPT models under that single subscription.

The Playbook

If your device can handle the localized inference, **Aider + Unsloth (running Hermes)** is the absolute best terminal setup. It gives you CLI control without the Claude Code price tag.

If you decide you *do* need cloud reasoning for a tough problem but refuse to pay token fees, download **Cursor** or use the **GitHub Copilot** extension. You pay the flat \$20 subscription, use their bundled Claude/GPT access, and cancel whenever.