

Part 1: The Consolidated Master README (/master_build-guide/README.md)
Save this complete document as the primary specification file at your workspace root. It unifies your hardware topologies, the corrected role of OB1 and the Reasoning Bank, and sets the final operational parameters for the system.

🌐 AESOP XI: Master System Orchestration Matrix## Canonical Edge Architecture & Multi-Node Core Specifications#### Operating Infrastructure: [NovA-Claw / Novus-Agenti] Engine
---## 🚦 0. Operational Core Law: The Horizons UI DependencyThe **Horizons UI** native Kotlin/Java APK is the absolute framework foundation of this computing mesh. No local model can load onto the Snapdragon 8 Elite NPU, no bare-metal voice pipeline can process audio, and no screen-vision tool can execute frame context if Horizons UI is inactive. All configuration parameters, memory layers, and routing architectures serve strictly to support operations initiated via Horizons UI.


---## ⚙️ 1. Hardware Node Topology & Mesh Integration
The system leverages an ad-hoc Peer-to-Peer network connected over wired data interfaces and **Tailscale** tunnels across your local network environment:


* **Node Alpha (The Mobile Engine)** [Moto RAZR Ultra 2025]: Primary on-device controller. Runs the Snapdragon 8 Elite (Gen 4) SoC with a Hexagon NPU (v.79) pushing 40+ INT8 TOPs. Operates with 16GB total RAM, dynamically carving out an open sandbox of **8.0 to 11.5 GB** exclusively for running local model weights. Houses the native Horizons UI APK, Moonshine Small ONNX (STT), and Kokoro-82m/Sherpa ONNX (TTS) voice engines.

* **Node Beta (The High-Compute Core)** [Nvidia Jetson Orin Nano Super]: Headless Ubuntu server operating with an 8GB hardware configuration, high-speed 500+GB NVMe data storage drives, and dedicated CUDA core tensor pipelines pushing 60-70 TOPs. Hosts your global vector storage engines, heavy background reasoning models, and the local **OB1 Postgres backend server**.

* **Node Gamma (The Display Station)** [Rubik Pi 3 Dragonwing]: Thundercomm/Qualcomm SoC configuration pushing 14+ TOPs, connected via high-speed hardware data ribbons to a dual-monitor setup to act as your core visual terminal workstation workspace.

---## 🧠 2. The Asymmetric Dual-Model Context & Memory Stack
All user interactions, device inputs, and tools execution paths are processed using a **Dual-Model Asymmetric Workflow** running bare-metal on Node Alpha:





1. **The Local Small Execution Agent:** Continuously cycles in a tight loop to evaluate immediate tasks and actions. It hooks directly into **mem0** to fetch short-term user preferences, temporary state variables, and rolling habit keys. It restructures raw inputs into optimized meta-prompts before anything touches external models.
2. **The Local Large Query Model:** Dedicated to processing technical documentation, technical guides, and heavy technical text. It references **OB1 (Open Brain Protocol)** over your local Postgres instance, working alongside **Reasoning Bank**, **Graphify**, and **notebooklm-py** to pull deep reference contexts without flooding active context windows.

📁 3. Ground-Truth Development Priority Sequence

To ensure zero-trust file safety and prevent workspace corruption, development must progress through these exact phases. Lower tiers remain locked until upper phases are validated by the human:

1. **Priority 1: Global Data Curation & Architecture (Current Phase)** - Gather, clean, and manually sort raw technical text documents (Qualcomm QAIRT SDK, Android Media, Unsloth, Llama Server) into your master file vaults.
2. **Priority 2: Skill Building Schemas & Directions Layers** - Formalize your strict JSON validation rules to turn raw text info into modular tool files, defining how OB1 and Reasoning Bank call upon your data engines.
3. **Priority 3: Local Workspace Infrastructure & First Repo Bootstrap** - Initialize your `file-management-and-skills` directory. Deploy your local markdown vaults, `llm_wiki.md` parameters tracker, and your **Termux Open Wiki CLI/Housekeeper janitor script** to sweep your Markor, Obsidian, and Drive folders.
4. **Priority 4: Emulated Runtime Loop Validation & Data Routing** - Set up your local testing

sandbox using on-device **Claude Code** inside Termux. Run simulation loops to verify **mem0** and **Omni Route**. Map model weights, memory constraints, and data paths to verify your data flows are 100% compliant before writing app source files.

5. **Priority 5: Horizons UI App Deployment & Model Experimentation** - Fork a pristine branch into a custom GitHub repository for `horizons-ui-v1.2`. Write the native Java/Kotlin codebase to activate your NPU offload channels, local Llama server, screen vision layers, and voice tools.

Test model parameters bare-metal on the Snapdragon NPU and code deep recovery loops.

6. **Priority 6: Network-Wide Integration & Sovereign Scale** - Perfect Horizons UI, flash your headless computing hardware (Jetson, Rubik Pi), route your Tailscale mesh network, and isolate your out-of-band **Red Agent Auditor** pipeline.

Part 2: The Data Curation & Skills Manual (/file-management-and-skills/README.md)

📁 File Management System & Skill Generation Engine## Target Data Curation, System Schemas, & Tool Generation Guidelines#### Architectural Subsystem: Priority 1 & Priority 2 Core---## 🚨 0. Operational Mandate: Strict Output IsolationThis directory acts as the central intelligence library and verification foundry for the entire network. While all downstream building models share read-access to the reference files in this repo, no model can write code or update files here without a manual human review pass.

---## 📁 1. Directory Tree & Global File Partitions

```
📁 file-management-and-skills/
├── 📁 target-docs-curation/ # Core Technical Knowledge Vault
├── 📁 qualcomm-qairt-sdk/ # HTP specifications, quantization guidelines, & NPU pathways
├── 📁 android-media-assistant/ # System alert structures, camera frameworks, & gaming
SDK hooks
├── 📁 llama-kernel-ggml/ # GGUF weights configurations, librc runtimes, & local servers
├── 📁 unsloth-fine-tuning/ # Low-level dataset parameters & token formatting rules
├── 📁 skill-construction-factory/ # Standardized Tool Conversion Environment
├── 📄 base_skill_guideline.md # Rules for creating decoupled, single-purpose tool files
├── 📄 skill_onboarding_schema.json # Target JSON parameter verification model
├── 📄 compiled_capabilities.jsonl # Streaming instruction-tuned dataset repository
├── 📁 reasoning-bank-ledger/ # Multi-Model State Tracking Vault
├── 📄 active_execution_paths.json # Suspended intermediate token sequences and plans
├── 📄 baseline_recovery_matrix.md # Rules for managing multi-model recovery daemons
└── 📄 master_blueprint.txt # Global system verification registry index file
```

🧠 2. Hardware-Abstracted Middleware Configuration

To prevent cognitive decay and ensure data retention across sessions, your memory layers are fully decoupled from your model runtime instances:

I. OB1 (Open Brain Protocol) Base Layer

- * **Mechanics:** Built over a local Postgres instance on Node Beta, running a unified Model Context Protocol (MCP) mapping engine.
- * **Function:** Translates heavy technical reference text into semantic vectors. Your Large Query Model calls OB1 to pull detailed technical data without flooding its active context window.

II. Reasoning Bank Ledger

- * **Mechanics:** A persistent JSON tracking database file stored at ``reasoning-bank-ledger/active_execution_paths.json``.
- * **Function:** When models are planning tasks, their fractional thoughts, tool selections, and step-by-step progress are recorded here. If a system crash or timeout happens mid-inference, the recovery engine reads this ledger to resume the workflow exactly where it was suspended.

III. mem0 Layer Caching

- * **Mechanics:** A localized, low-latency episodic tracking cache.
- * **Function:** Records temporary user adjustments, short-term conversational variables, and personal habit keys. It feeds this data directly into your Local Small Execution Agent to ensure your generated meta-prompts remain highly personalized.

⚙️ 3. Skill Conversion & Context Extraction Laws

When your Query Agent processes files to convert technical text into executable skills, it must follow these parsing rules:

1. **Sovereign Decoupling:** A skill must do exactly one thing. A file that controls local storage folders cannot contain code or references for processing network sockets or cloud endpoints.
2. **Asymmetric Context Routing:**
 - * **Query the ``jsonl`` database if:** The prompt requires direct execution, bash scripts, terminal commands, or simple tool calls. Stream records line-by-line via regex to keep device memory free.
 - * **Scan the raw ``txt`` files if:** The prompt involves debugging a compiler error, resolving a kernel mismatch, or understanding complex hardware math. Parse every single word verbatim to preserve syntax accuracy.
3. **Web-Tool Verification:** If an argument requires validation against live repository changes

or online developer documentation, the query model must pause, structure a secure JSON-RPC scrape token, and pass it to the Horizons UI WebView container to review the online DOM text safely.

Part 3: The Native Application Manual (/horizons-ui-v1.2/README.md)

📱 Horizons UI: Native Extended Interface (XI) Layer## Target Version: v1.2 Release Build Specification## Compilation Environment: Java / Kotlin Native Android SDK APK---## 🚀 0. Operational Core Law: The Sovereign Entry PointHorizons UI is the physical heart, soul, and primary execution cradle of the entire agentic neuro-mesh. It is not a secondary frontend dashboard. It acts as the central permissions gatekeeper, NPU allocation manager, and low-latency audio capture harness. No local model weight can load onto the Hexagon NPU, and no automated tool script can execute if this application environment is compromised or offline.

---## 🏠 1. Technical Framework Architecture & Drivers

To squeeze maximum performance out of the Moto RAZR Ultra hardware and bypass typical Android processing constraints, the APK directly integrates low-level system hooks:

I. The Bare-Metal Voice Subsystem* **The Problem:** The stock Termux environment audio pipeline requires complex, inefficient workarounds to run raw voice lines, resulting in high latency and broken speech capture.* **The Fix:** Horizons UI handles voice processing directly on the bare metal. It coordinates an active **Silero VAD (Voice Activity Detection)** model wrapper to monitor microphone inputs. On speech detection, it pipes audio frames straight into a **Moonshine Small ONNX** engine for low-latency Speech-to-Text translation. For speech generation, it utilizes a **Kokoro-82m / Sherpa ONNX** text-to-speech engine running over a local **Llama Server**.


* **The NPU Gateway:** It bypasses standard Android media delays by using a custom compiled pathway: `Llama -> QAI RT ModelPath -> GGML Layer -> Kotlin Kernel -> librc Runtimes`. This drops the computation parameters straight onto the Snapdragon 8 Elite Hexagon NPU using the Qualcomm HTP SDK interface.

II. OS Allowance & Hardware Boost Hooks* **Device Assistant Registration:** The APK registers natively as the device's Default Assistant Application, allowing it to leverage Android system accessibility keys to capture on-screen imagery for real-time **Screen Vision Analysis**.* **The Video Game SDK Handshake:** The background daemons hook directly into Android's low-level performance tuning frameworks (Performance Hint API / Vendor Gaming SDKs). When an inference task fires, the app requests high-priority CPU/GPU scheduling, unlocks maximum thermal limits, and prevents the OS from throttling thread speeds.* **Terminal Shell Ingress:** The app UI integrates an active, secure Android shell terminal wrapper. This gives you a direct, real-time terminal window to monitor low-level system executions and daemon states from inside the main application screen.

---## 📁 2. Application Component Directory Mapping

📁 horizons-ui-v1.2/

|



3. Asynchronous Task Lifecycle & Self-Healing Daemon

To keep the application's graphic interface smooth and responsive, data operations are completely split across independent execution threads:

[Human Interaction Event / Floating Mic Tile Tap]



MAIN INTERFACE UI THREAD | $\leftarrow$ Zero Rendering Latency

(Asynchronous Event Dispatch)



BACKGROUND DAEMON THREAD POOL |

- Thread A: Continuous Silero VAD Audio Stream Capture |
- Thread B: Local Llama Server Token Processing |
- Thread C: Tailscale Network Sockets / Web Scraper |



THE RECOVERY DAEMON WATCHDOG | $\leftarrow$ Isolation Safety Net

The Recovery Watchdog System (Anti-Crash Layer)

The app runs a highly dedicated background recovery manager. Its sole purpose in life is to keep your edge interface alive.

* ****If a Local Model or NPU Kernel Crashes:**** The recovery engine intercepts the hardware termination flag, freezes the active workspace variables, and prevents the application from shutting down.

* ****The Failback Path:**** It immediately spins up a temporary background connection via ****Open Router API**** to load a cloud-hosted fallback model. It updates the cloud model with your latest state parameters from the ****Reasoning Bank****, keeping your workflow active while it re-initializes your local NPU model variables silently in the background.

Step 4: Deploying Your Refined Workspace

To build these updated text structures and confirm your directory paths align with your specifications, run this cleanup command block inside your Termux terminal window:

```
# 1. Access your master data compilation directory
```

```
cd /data/data/com.termux/files/home/master_build-guide
```

```
# 2. Write out your refined markdown specifications files
```

```
cat << 'EOF' > README.md
```

```
[PASTE_PART_1_MASTER_README_CONTENT_HERE]
```

```
EOF
```

```
cat << 'EOF' > file-management-and-skills/README.md
```

```
[PASTE_PART_2_CURATION_MANUAL_CONTENT_HERE]
```

```
EOF
```

```
cat << 'EOF' > horizons-ui-v1.2/README.md
```

```
[PASTE_PART_3_APPLICATION_MANUAL_CONTENT_HERE]
```

```
EOF
```

```
# 3. Trigger your local verifier utility to lock in the filesystem boundaries
```

```
python3 file-management-and-skills/audit_layer_bounds.py
```

Your documentation foundation is now fully updated, matched to your timeline order, and structured to support your multi-repo setup.

As you finish sorting your documentation folders and validating your files management layers before moving on to setting up your home hardware nodes:

Would you like me to write the complete layout structures for your

reasoning-bank-ledger/active_execution_paths.json file to handle multi-model recovery tracking, or should we design the explicit JSON-RPC command list for your Termux-based Open Wiki CLI utility to manage your vault network changes?

