

The ecosystem has evolved through multiple naming conventions and target iterations—ranging from early terminal-native sandboxes (**AESOP**, **Termux CLI Stack**) to native on-device hardware engines (**Novus-Agenti**, **NovA-Claw**, **N0.V4**, and the upgraded **Horizons Workbench UI**).

The canonical repository blueprint requires sorting everything into **Four Core Buckets** to prevent your engineering agents from drifting into old contexts or hallucinating stale runtime targets.

Tier 1: The Cold Archive & Source References (/raw_sources/)

These are heavy, static background documents, technical metrics, and community tutorials that the model should use strictly as read-only reference data—never as active working rules.

- **Stanford AI Index & General Trends:**
 - Artificial Intelligence Index Report | Stanford HAI (Massive industry benchmark data bank)
- **External Technical / Framework Discoveries:**
 - MEMO_ A Modular Framework for Training a Dedicated Memory Model...pdf (Parametric memory layer research)
 - Together AI Open-Sources OSCAR...pdf (2-bit attention-aware KV cache optimization spec)
 - NVIDIA Releases Polar...pdf (Gateway rollout framework for token-faithful GRPO training loops)
 - Meet GitAgent..., OpenAI Releases Symphony..., How CopilotKit Is Redefining..., Meet GitHub Spec-Kit... (Framework mappings for agentic orchestration)
 - Meet OmniVoice Studio...pdf (Tauri/FastAPI local alternative reference stack)
 - Hexo Labs Open-Sources SIA... (Self-improving agent scaffold loops)
- **Qualcomm / Hardware Specific Specs:**
 - qualcomm/Phi-4-Mini-Instruct · Hugging Face, Phi-4-Mini-Instruct - Qualcomm AI Hub
 - qualcomm/Qwen3.5-0.8B · Hugging Face, Qwen3.5-0.8B - Qualcomm AI Hub, unsloth/Qwen3.5-9B-GGUF · Hugging Face
 - Gemma-4-E4B-it - Qualcomm AI Hub
 - qualcomm/LiteHRNet · Hugging Face, qualcomm/PiperTTS-DE · Hugging Face
 - Liquid AI Releases LFM2.5-8B-A1B...
 - 8.2mb what is a hybridized agentic on device neuro mesh...pdf
- **Keyboard & Input Utilities:**
 - FUTO Keyboard, Voice Input Models - FUTO Keyboard

Tier 2: Core Architecture Vault (/master_wiki/)

This is the single source of truth for design constraints, structural layouts, and evolution reports mapping out the native system topology.

- **Native Android (Kotlin/Compose) Next-Gen Workbenches:**
 - HORIZONS_state_and_corrections_MASTER.md (**Canonical State Document** detailing structural panel layouts, model layer bindings, and the credential list rewrite)
 - ARCHITECTURE_REVIEW_v3_amendment.md (Crucial dual-axis framework detailing the

Cloud Failover Ladder vs. the VLM Context Dispatcher)

- The Horizons Workbench: A Guide to Your Modular Intelligence Engine
- 4. Technical Specification: The Four Rooms & Seven-Tile Modular Architecture
- The Horizons Workflow: From Storage to Switch
- HORIZONS_UI_SPEC_v3.md (Detailed single-Activity layout and Accessibility Service automation blueprints)
- NO_V4_ARCHITECTURE_v3.md (Distributed edge-cloud specification targeting Snapdragon 8 Elite hardware metrics)
- Architecting the Agentic Neural Mesh for Snapdragon 8 Elite, Agentic Neural Mesh: Master Architecture and Evolution Report
- **Legacy / Interim Sandbox Ecosystem Templates (Marked as Contextual Reference):**
 - TERMUX_CLI_STACK.md (Bridges the gap; maps out the native headless terminal environment layer utilized by RUN_COMMAND)
 - Copy of SOTU-2026-07-27.md (Historical project state tracking across legacy variations)
 - deploy_T3-BRINGUP.md (Infrastructure specifications mapping the Tier 3 home network hub target: Jetson Orin Nano Super + Rubik Pi 3 desktop-OS node over Tailscale)

Tier 3: The Almanac & Skill Definitions (/inference_skills/)

Line-delimited facts, structured JSONL logs, terminal execution setups, and open SKILL.md plugin files used to instantly direct local engines.

- **Active Agent Tool Routing Rules:**
 - horizons-wiki-SKILL.md (Packages stable code-review architecture, daemon path parameters, and Anthropic ephemeral caching boundaries into a system block prefix)
 - 1b How to Structure Your Wiki Folders..., Turning a document into a Skill by adding Frontmatter... (Operational instructions for compiling modular .txt or .json directories into queryable skill blocks)
 - claude-skills.md (Reference domain bundle for external tool building/auditing frameworks)
- **Terminal Environment & API Wiring Recipes:**
 - GENIEX-DAEMON-PLAN.md / Copy of GENIEX-DAEMON-PLAN.md (Blueprints for launching standard OpenAI request formats via a standalone geniex serve native process over port 18181)
 - Standardized UNIX socket protocol, (Zero-copy IPC optimization blueprints utilizing shared memory fields)
 - llm-wiki-termux-setup-full-notes.md / llm-wiki-termux-setup-urls.txt (Headless environment installation logic)

Tier 4: Session Execution & Hygiene Logs (/automation_scripts/)

Active working tracks, debugging sequences, failure logs, and step-by-step master checklists used to eliminate sessional drift.

- **Current Iteration Logs:**
 - BUILD-ACTION-PLAN.md / Copy of BUILD-ACTION-PLAN.md (Parallel processing partitioning guidelines for Repo vs. Device workflows)
 - FAILURE_LOG_SESSION_20260527.md / weekly_audit_prompt.md (Failure frequency

- tracker monitoring degradation types like specification drift or tool error loops)
- NeuroOmni_agy_build_pack.md (Milestone execution checkpoints for scaffolding the base project architecture)

Step-by-Step Execution for Unified Analysis

To prevent context corruption, do not load heavy references simultaneously with the active target code. Execute your unilateral analysis using the following pipeline:

1. **Lock the Base Pool:** Use **HORIZONS_state_and_corrections_MASTER.md** combined with **ARCHITECTURE_REVIEW_v3_amendment.md** as your absolute architectural anchor.
2. **Evaluate Path Congruence:** Use **BUILD-ACTION-PLAN.md** to verify that file alterations stay completely distinct across active sub-sessions.
3. **Run the Audit Matrix:** Filter incoming tracking data using the failure categories detailed in **weekly_audit_prompt.md** to verify execution readiness before generating patches.