


To orchestrate multiple frontends and pluggable backends while enforcing **QAIRT (Qualcomm AI Runtime)** for your GGUF asset pipeline, you need to design a **Hexagon-Optimized Local Orchestrator Pattern**.

Since QAIRT cannot execute raw GGUF files dynamically, your orchestrator must handle a **dual-payload asset pattern** behind the scenes. It will map your user-facing .gguf file to its corresponding pre-compiled .bin context binary when the QAIRT backend is active, while allowing other backends (like llama.cpp or ONNX Runtime) to ingest the raw .gguf directly when swapped in.

Here is the architectural blueprint to build this pluggable orchestration engine in Kotlin and Java.

1. The Core Architecture Diagram

Your orchestrator sits as a clean abstraction layer. Neither your UI frontends nor your raw model files talk directly to the hardware.



2. Define the Pluggable Backend Interface

Create an explicit, lifecycle-managed backend contract. This allows you to hot-swap engines at

runtime without forcing your UI frontends to recompile or change their state observers.

```
// InferenceBackend.kt
interface InferenceBackend {
    val backendId: String
    val isHardwareAvailable: Boolean

    // Pass execution configurations dynamically
    fun initialize(modelPath: String, context: Context): Boolean
    fun generateTokenStream(prompt: String): Flow<String>
    fun release()
}
```

3. Implement the QAIRT Pluggable Engine

The QAIRT implementation requires a paired-file lookup mapping strategy. When your orchestrator tells it to load qwen3.5.gguf, it intercepts the path and loads the hardware-compiled .bin target.

```
// QairtHtpBackend.kt
import com.qualcomm.qairt.runtime.QnnBackend
import com.qualcomm.qairt.runtime.QnnContext

class QairtHtpBackend : InferenceBackend {
    override val backendId = "QUALCOMM_QAIRT_HTTP"
    private var qnnContext: QnnContext? = null

    override val isHardwareAvailable: Boolean
    get() = try {
        // Check if the physical HTP engine chips are active
        val backend = QnnBackend("libQnnHtp.so")
        backend.isBackendSupported(QnnBackend.Property.HTP_AVAILABLE)
    } catch (e: Exception) {
        false
    }

    override fun initialize(modelPath: String, context: Context): Boolean {
        return try {
            // Orchestration Trick: Map the requested .gguf to the QAIRT binary asset
            val binaryPath = modelPath.replace(".gguf", "_http.bin")
            val modelBytes = context.assets.open(binaryPath).readBytes()
        } catch (e: Exception) {
            false
        }
    }
}
```

```

        val nativeBackend = QnnBackend(context.applicationInfo.nativeLibraryDir +
"/libQnnHtp.so")
        qnnContext = nativeBackend.createContext(modelBytes)
        true
    } catch (e: Exception) {
        Log.e("QAIRT_BACKEND", "Failed to bind GGUF asset to QAIRT engine", e)
        false
    }
}

override fun generateTokenStream(prompt: String): Flow<String> = flow {
    // Core execution loop utilizing native QAIRT graph inputs/outputs
    // Emit strings token-by-token down the flow pipeline
}

override fun release() {
    qnnContext?.close()
}
}

```

4. Build the Central Master Orchestrator

The central orchestrator manages the registry of your backends, detects the device's hardware capabilities, hooks into multiple UI frontends via standard architectures, and implements a automatic fallback chain.

```

// ModelOrchestrator.kt
import kotlinx.coroutines.flow.MutableStateFlow
import kotlinx.coroutines.flow.StateFlow

class ModelOrchestrator(private val context: Context) {

    private val backendRegistry = mutableMapOf<String, InferenceBackend>()

    private val _activeBackend = MutableStateFlow<InferenceBackend?>(null)
    val activeBackend: StateFlow<InferenceBackend?> = _activeBackend

    init {
        // Register all pluggable engines dynamically
        registerBackend(QairtHtpBackend())
        registerBackend(LlamaCppHexagonBackend()) // Secondary high-speed backup
        registerBackend(FallbackCpuBackend()) // Safety baseline
    }
}

```

```

    }

    fun registerBackend(backend: InferenceBackend) {
        backendRegistry[backend.backendId] = backend
    }

    // Frontends invoke this to hot-swap engines seamlessly mid-session
    fun switchBackend(backendId: String, modelName: String): Boolean {
        val targetBackend = backendRegistry[backendId] ?: return false

        if (!targetBackend.isHardwareAvailable) {
            Log.w("ORCHESTRATOR", "Requested hardware for $backendId missing. Skipping.")
            return false
        }

        _activeBackend.value?.release()

        val success = targetBackend.initialize(modelName, context)
        if (success) {
            _activeBackend.value = targetBackend
        }
        return success
    }

    // Automatic hardware matching orchestration routing
    fun autoSelectBestBackend(modelName: String) {
        when {
            backendRegistry["QUALCOMM_QAIRT_HTP"]?.isHardwareAvailable == true -> {
                switchBackend("QUALCOMM_QAIRT_HTP", modelName)
            }
            backendRegistry["LLAMACPP_HEXAGON"]?.isHardwareAvailable == true -> {
                switchBackend("LLAMACPP_HEXAGON", modelName)
            }
            else -> {
                switchBackend("FALLBACK_CPU", modelName)
            }
        }
    }
}

```

5. Hooking up Multiple Frontends (Pluggable UI Layer)

Because your orchestrator surfaces model generation as a standard reactive streaming construct (`Flow<String>`), any type of frontend interface can register to receive data pipelines

concurrently without knowing which backend engine is running.

Frontend Variant A: Jetpack Compose UI

```
@Composable
fun ChatScreen(orchestrator: ModelOrchestrator) {
    val activeEngine by orchestrator.activeBackend.collectAsState()
    var chatTokens by remember { mutableStateOf("") }

    Button(onClick = {
        coroutineScope.launch {
            activeEngine?.generateTokenStream("Hello Qwen!").collect { token ->
                chatTokens += token
            }
        }
    }) {
        Text("Run via ${activeEngine?.backendId ?: "None"}")
    }
}
```

Frontend Variant B: Background Service / IPC / Headless Terminal

```
class HeadlessAILayerService : Service() {
    private lateinit var orchestrator: ModelOrchestrator

    // Accessible over AIDL, BroadcastReceivers, or Local Sockets
    fun handleSystemRequest(prompt: String) {
        CoroutineScope(Dispatchers.Default).launch {
            orchestrator.activeBackend.value?.generateTokenStream(prompt)?.collect { token
->
                sendBroadcastToSubsystems(token)
            }
        }
    }
}
```

6. Managing the GitHub CI Artifact Matrix

To ensure your orchestrator has the files it needs for **both** the plug-and-play GGUF engines and QAI RT HTP, your CI pipeline needs to export a unified directory structure inside your APK asset

space.

Configure your GitHub Actions artifact compiler to output matching target names:

```
# Inside .github/workflows/compile-assets.yml
- name: Package Orchestrator Asset Map
  run: |
    # Move raw weights for fallback and native GGUF backends
    cp qwen3.5-9b-q4_0.gguf ./app/src/main/assets/models/qwen3.5.gguf

    # Run the QAIRT toolchain to make the matching hardware partner binary
    python3 $QNN_SDK_ROOT/bin/gguf_builder \
      --input_model qwen3.5-9b-q4_0.gguf \
      --target_backend htp \
      --output_path ./app/src/main/assets/models/qwen3.5_htp.bin
```

Summary Orchestrator Matrix

Component Layer	Responsibility	Structural Strategy
Interface Definition	Uniform parameter mapping.	Keeps frontends clean of C++ dependencies.
QAIRT Backend	Intercepts paths dynamically.	Looks up .bin counterparts silently while claiming to load .gguf.
Central Registry	Runtime engine hot-swapping.	Enables fallbacks if specific Snapdragon chips lack space/support.

If you'd like, let me know:

- How you intend to pass prompts between your frontends and the orchestrator (such as **Kotlin Channels**, **Shared Flows**, or **AIDL Binder interfaces**).
- Whether you need a specific **RAM management strategy** to evict models cleanly when hot-swapping backends on memory-constrained devices.

I can provide the exact implementation details for memory-safe model unloading or cross-process IPC communication!