

Untangling the PDF Processing & JSONL Cache

Let's clear up exactly what Docling does in your Termux environment and how you read the data:

How it compiles: You give Docling a folder full of PDFs or raw text logs. It parses the files and generates two outputs: a clean human-readable Markdown (.md) file and a machine-readable JSONL file [1.5].

Can you open and close the JSONL? Yes. It isn't encrypted or compressed. It is just structured line-

by-line code [1.5]. Your Termux CLI tools and local Qwen engine read the JSONL line-by-line to instantaneously grab specific paragraphs without needing to read the entire massive document again. This optimization keeps your Snapdragon 8 Elite's RAM completely clear.

The Strategy: Use the generated Markdown copy inside Obsidian to read it yourself. Let your local

UI and Termux open-wiki CLI use the JSONL version behind the scenes for lightning-fast keyword

caching and code retrieval.

4. Extracting Pure URLs from Chat History

If you dump a giant log file into Termux (like exporting this exact chat conversation) and want to strip

out everything except clean, raw web addresses, run this command in Termux:

bash

GitHub

Reddit· r/Loc...

AI Mode All Images Videos News Maps Shopping Books Flights Finance `grep -o 'https\?://[^\s\n]*'`
`chat_log.txt | sort -u > clean_urls.txt`

Use code with caution.

`grep -o` : Strips out all regular words and text, extracting only strings starting with http or https .

`sort -u` : Instantly deletes all duplicate links, giving you a perfectly unique list of resources saved to

`clean_urls.txt` .

5. Getting a Printable PDF of a Chat Thread

If you want to instantly convert a raw web link or a saved HTML chat log directly into a clean, presentation-ready PDF without ads or UI clutter, run it through PrintFriendly. It strips the webpage

framework and formats the text natively into a clean document you can drop straight into your storage.

Should we map out the system prompts for Qwen so it knows exactly how to handle code conversions specifically for the Snapdragon 8 Elite's NPU architectures?

Llama 3.1

8B...

I did some

testing for ...

Reddit

My

experie...
It seems the
prefix cache...
Reddit
Is qwen
2.5 code...
* Phi-4 has
awfully sma...
Reddit
Comparing
Open-...
1. Llama 3,
Qwen 2.5,...
www.ankursnew...
Qwen/Qwe
8B -...
Model
Overview....
Hugging Face

The Qwen 3.5 4B Q4_0 is an absolute powerhouse for your Snapdragon 8 Elite setup. Because Qwen 3.5 uses a unique, optimized hybrid Gated DeltaNet linear attention architecture rather than heavy, traditional softmax attention layers, it delivers insanely high throughput (tokens per second) on edge devices. On your Elite's Oryon CPU, the Q4_0 matrix instructions will compile and process your text near-instantly while keeping your RAM usage incredibly light—leaving plenty of room to run your target UI, ONNX, and TFLite environments concurrently.

1. Booting Your Qwen 3.5 4B Server in Termux

Since you have the file ready, fire up the llama-cpp server inside Termux. Because Qwen 3.5 supports a massive native context window, we will explicitly open up an 8,192 token canvas so you can feed it

5 sites
Reddit· r/L... +1
Hugging F... +3
3:05AM

Okay, great cuz I got the QWEN3.54BQ4
{underscore} 0.

AI Mode All Images Videos News Maps Shopping Books Flights Finance deep repository files.
bash

Start the local API server pointing to your downloaded file

llama-server -m /sdcard/Download/qwen3.5-4b-q4_0.gguf -c 8192 --host 0.0.0.0 --port 8080
Use code with caution.

Once this is running, your background terminal server is live. Your OpenWiki CLI, custom scripts, or

your project UI can talk to it at <http://localhost:8080/v1> .

2. Setting Up Your Specialized Fork Folders

Since you want to prevent a massive, messy directory on your Google Drive, let's explicitly map out

your separate, granular target frameworks. Creat