

AESOP XI Architecture Blueprint Update

<https://gemini.google.com/app/9c16534ff86ba138>

User prompt: The first file is the architecture template it's beautifully written but full of old information can you take these attached more recent documents and implement all of these updates and make it a new copy I want to save the original but let's make another a new version of this architecture document updated with all this new information

Response: An updated master architecture document integrates the 3-APK Android topology, the dual-tier personal-vs-business GCP Vertex AI Agent Builder credit routing, the RLVR continuous fine-tuning data flywheel, the expanded tool/skill harnesses, and the on-device inference orchestration loop.



AESOP XI: Autonomous Edge-Computing Architecture Blueprint

Heterogeneous System Orchestration, Asymmetric Routing, & Autonomous Agent Flywheels

Specification Version: v2.0-Production-Ready | Target: Multi-Node Sovereign Mesh

1. Canonical Build Sequence & Roadmap

Phase 1: Knowledge Curation & Universal JSONL Marker Indexing
Phase 2: Tool Harness Ingestion (Pocock, Honey for Devs, Prime Agent, ECC, GSD)
Phase 3: Repository Layout & Multi-Agent Grill Session Bootstrapping
Phase 4: 3-APK Engine Development (Horizons UI + Shell Daemon + Media Daemon)
Phase 5: Local Dual-Model Inference & Continuous Improvement RLVR Loop (GCP)
Phase 6: Multi-Node Mesh Orchestration & Autonomous Agent Swarm (Jetson / Pi)


- **Priority 1: Knowledge Curation & Universal JSONL Marker Indexing**
 - Ingest and structure technical documentation: Qualcomm QAIRT/HTP SDK, Android Media/Accessibility APIs, Google Vertex AI Agent Builder, Llama Server, and Unsloth.
 - Attach dense JSONL marker schemes to every markdown file, skill directory, and tool repository to allow instant KAG/RAG indexing without context window bloat.
- **Priority 2: Skill Building Schemas & Tool Harness Integration**
 - Transform static reference repositories into machine-executable skills.
 - Ingest developer harnesses: **Matt Pocock Skills**, **Honey for Devs**, **Prime Agent** (RLM persistent REPL + Continual Harness with rollback snapshots), **ECC**, **GSD (Get Shit Done)**, **Claude Video**, **Reverse Skills**, **Code Review Graph**, **mem0**, and **OpenAI/Anthropic LocalAI Omni-Route bridges**.
- **Priority 3: Repository Bootstrapping & Architectural Grill Session**
 - Deploy `file-management-and-skills` as the master root blueprint.
 - Execute the "Grill with Docs" orchestration agent to establish repository trees, assign individual agent manifests, and manage cross-repo directory boundaries.
- **Priority 4: 3-APK Architecture Deployment (Parallel with Priority 5)**
 - Build the decoupled, multi-APK Android system: **Horizons UI** (Orchestrator), **Shell Daemon APK** (Termux-free OS assistant/accessibility access), and **Media Daemon APK** (Screen-vision + Silero VAD + Speech stack).
- **Priority 5: Dual-Model On-Device Inference & Concierge Gateway**
 - Deploy the dual-model tandem: **Qwen 3.5 0.8B GGUF Q4_0** (Executor) and **Qwen 3.5 9B GGUF Q4_0** (Query / Metaprompt compiler).
 - Wire the on-device **Concierge Flow**: Voice capture → Silero VAD → Screen Vision context → Synthesized clean markdown meta-prompt → User verification → Frontier Model / LocalAI → Kokoro TTS speech output.
- **Priority 6: Network-Wide Integration & GCP Cloud Flywheel**
 - Establish the continuous improvement training flywheel: Verified task-verdict triples collected locally → Exported to GCS → GCP Vertex AI / Compute Instances run RLVR (GRPO/Unsloth) fine-tuning → Updated weights pushed back to edge.
 - Activate the full multi-agent home node: **Red Agent Auditor**, **Housekeeping/Logs Compiler**, **Help Desk Manual Agent**, and **NPU Inference Manager**.

2. Multi-Node Compute Infrastructure


Node ID	Physical Platform	Primary Function	Core Acceleration & Compute	Primary Software & Daemons
NODE ALPHA	Moto Razr Ultra 2025	Local Orchestration & Concierge Hub	Snapdragon 8 Elite (Hexagon NPU v79)	3-APK Unified Stack (Horizons + Daemons), LocalAI, Moonshine/Kokoro
NODE BETA	Nvidia Jetson Orin Nano Super	Headless Home Server & Vector Hub	Dedicated CUDA Core Array (60-70 TOPs)	Postgres OB1 Protocol, Red Agent Auditor, Multi-Agent Swarm
NODE GAMMA	Rubik Pi 3 (Dragonwing)	Display Workstation & Helpdesk Gateway	Thundercomm / Qualcomm SoC (14+ TOPs)	Dual-Monitor Display Engine, System Log Viewer
NODE DELTA	Google Cloud Platform	Asymmetric RAG & Heavy RLVR Training	Cloud TPUs / A100/H100 Instances	Vertex AI Agent Builder (\$1,000 Credit Tier), GCS Buckets, Unsloth/Axolotl GRPO


3. The 3-APK Unified Android Architecture

Instead of running inside a restricted sandboxed terminal environment, on-device operations are split across three native APK layers:



4. End-to-End System Dataflow & Agent Concierge Pipeline





5. Multi-Agent Home Node Topology

The home node runtime (Node Beta / Jetson Orin) hosts a specialized team of autonomous sub-agents:

- **Red Agent Auditor & Safety Guardrail:** Intercepts outgoing model execution payloads, filters execution traces against the Nope Data Bank, and signs off on local .jsonl telemetry logs.
- **Home Node Housekeeper & Log Compiler:** Aggregates execution logs across nodes, dedupes trace histories, manages disk allocations, and formats training data pairs.
- **IT Help Desk & Manual Operator Agent:** Searches local technical manuals, tool specifications, and hardware registries to debug operational failures.
- **Web Ingestion & News Monitor Agent:** Tracks breaking updates in local open-weight models, new tools, and upstream repository changes.
- **NPU / Local Inference Manager:** Balances weight offloading between Snapdragon Hexagon NPU, Jetson CUDA cores, and fallback cloud APIs.

6. GCP Cross-Account IAM Handshake & RLVR Loop

To utilize developer credits without duplicating proprietary data, cloud infrastructure is split into an asymmetric two-tier architecture:

- **Business / Contractor Resource Tier:** Holds source documentation in gs://business-secure-vault-bucket/. Grants read-only roles/storage.objectViewer to the personal consumer service account.
- **Personal Developer Tier (\$1,000 Credit Pool):** Vertex AI Discovery Engine (service-PERSONAL_PROJECT_NUMBER@gcp-sa-discoveryengine.iam.gserviceaccount.com) points directly to the business bucket without copying data, absorbing 100% of RAG query costs.
- **The RLVR Training Loop:** Local task-output-verdict triples stream to GCS. On a scheduled trigger, Vertex AI / GCP Compute instances execute a GRPO fine-tuning run on fixed base weights (Qwen/Llama) using Unsloth/Axolotl, deploying tuned GGUF weights back down to Node Alpha and Beta.

7. Decentralized Multi-Repository Map

```

📁 master_workspace/
├── 📁 horizons-ui-v2.0/           # Kotlin/Java APK 1: Orchestration GUI, Webview, Chat Tile[cite: 8, 9]
├── 📁 nova-daemon-shell/         # Kotlin/Java APK 2: OS Assistant & Accessibility Sandbox Bypass[cite: 8,
├── 📁 nova-daemon-media/         # Kotlin/Java APK 3: Screen Vision, Moonshine STT, Kokoro TTS[cite: 1, 8,
├── 📁 file-management-and-skills/ # Master Canonical Docs, JSONL Markers, Tool Schemas[cite: 1, 4]
├── 📁 agent-harness-hub/         # Prime Agent (RLM), Pocock Skills, Honey for Devs, ECC, GSD[cite: 1, 7]
├── 📁 localai-omni-route/        # Modality Routing, OpenAI/Anthropic Translation Endpoints[cite: 1, 6]
├── 📁 obsidian-knowledge-vault/   # Graphify Networks, KAG/RAG Markdown Trees, NotebookLM Hooks[cite: 1]
├── 📁 red-agent-auditor/         # Air-gapped Execution Verification, Nope Data Bank[cite: 1]
├── 📁 node-beta-jetson/          # Postgres OB1 Vector Store, CUDA Inference Servers[cite: 1]
└── 📁 gcp-training-flywheel/     # RLVR Export Scripts, Cross-Account IAM Handshake, GRPO Pipelines[cite: 1]
  
```

8. Universal JSONL Marker Schema

Every markdown documentation file and skill repository includes an attached .jsonl schema block to enable rapid, low-token search indexing across local agents and Vertex AI Discovery engines:

```
{
  "record_id": "SKILL_POCKOCK_GRILL_001",
  "document_path": "agent-harness-hub/pockock-skills/grill-me.md",
  "category": "SKILL_HARNESS",
  "target_runtime": "PRIME_AGENT_RLM",
  "metadata": {
    "title": "Grill-Me Architectural Interview Harness",
    "description": "Interactive question-driven prompt for technical planning",
    "primary_tools": ["claude-code", "open-wiki-cli", "prime-agent"],
    "required_context_keys": ["repo_scope", "target_architecture"]
  },
  "retrieval_tokens": ["grill", "spec", "architecture", "audit", "pockock"],
  "entry_points": {
    "repl_command": "/skill run grill-me",
    "jsonrpc_method": "agent.skills.execute"
  }
}
```

User prompt: Yeah that looks pretty good except the red agent auditor only intercepts when the context is loaded enough for actual training compilation it's the home assistant agent that does the cross agent auditing when where I'm running inference between all my models and that's going to be the one that's doing the daily log doing the initial script compilation that's going to actually get the thing that gets audited. There's also other little details that are either left out or not his current like the exact protocols and agent set up said I'm going to be running various applications and whatnot. I'm also going to be running grill with dogs because I've just spent a bunch of time getting my baseline documents all sorted out

Response: AESOP XI: Autonomous Edge-Computing Architecture Blueprint


Heterogeneous System Orchestration, Asymmetric Routing, & Autonomous Agent Flywheels


Specification Version: v2.1-Production-Ready | Target: Multi-Node Sovereign Mesh

1. Canonical Build Sequence & Execution Roadmap

Phase 1: Knowledge Curation & Universal JSONL Marker Indexing
 Phase 2: Tool Harness Ingestion (Pockock, Honey for Devs, Prime Agent, ECC, GSD)
 Phase 3: Grill with Docs Master Session & Repo Partitioning
 Phase 4: 3-APK Engine Deployment (Horizons UI + Shell Daemon + Media Daemon)
 Phase 5: Local Dual-Model Inference & Continuous Improvement RLVR Loop (GCP)
 Phase 6: Multi-Node Mesh Orchestration & Autonomous Agent Swarm (Jetson / Pi)

- **Priority 1: Global Knowledge Curation & Universal JSONL Indexing**
 - Assemble, clean, and partition core technical documentation: Qualcomm QAI/RT/HTP SDK, Android Media/Accessibility APIs, Google Vertex AI Agent Builder, Llama Server, and Unsloth.
 - Attach dense JSONL marker schemes to every markdown file, tool repository, and skill guideline to allow direct KAG/RAG indexing without context bloat.
- **Priority 2: Skill Building Schemas & Tool Harness Ingestion**
 - Transform static reference repositories into machine-executable skills.
 - Ingest developer harnesses: **Matt Pockock Skills** (Grill-Me, Spec, Ticket, TDD, Code-Review), **Honey for Devs**, **Prime Agent** (RLM persistent REPL + Continual Harness with rollback snapshots), **ECC**, **GSD (Get Shit Done)**, **Claude Video**, **Reverse Skills**, **Code Review Graph**, **mem0**, and **OpenAI/Anthropic LocalAI Omni-Route bridges**.
- **Priority 3: The "Grill with Docs" Master Architectural Session**
 - Deploy the `file-management-and-skills` repository as the canonical ground-truth baseline.
 - Run the interactive "Grill with Docs" orchestration agent to interrogate baseline documentation, assign discrete agent manifests, establish strict repository isolation boundaries, and roadmap tool utility across all harnesses.
- **Priority 4: 3-APK Native Android Deployment (Parallel with Priority 5)**
 - Deploy the decoupled 3-APK Android execution suite: **Horizons UI** (Orchestrator), **Shell Daemon APK** (Termux-free OS assistant/accessibility access), and **Media Daemon APK** (Screen-vision + Silero VAD + Speech stack).
- **Priority 5: Dual-Model On-Device Inference & Concierge Gateway**
 - Deploy the dual-model tandem: **Qwen 3.5 0.8B GGUF Q4_0** (Executor) and **Qwen 3.5 9B GGUF Q4_0** (Query / Metaprompt compiler).





5. Multi-Agent Home Node Topology (Node Beta / Jetson)

- **Home Assistant (Cross-Agent Auditor & Housekeeper):** Manages real-time cross-agent auditing during live inference sessions, aggregates daily execution logs, compiles draft automation scripts, and packages raw telemetry pairs for downstream training.
- **Red Agent Auditor (Batch Training Guardrail):** Operates strictly when accumulated context reaches the training compilation threshold, conducting air-gapped safety verification against the Nope Data Bank before datasets are shipped to GCP.
- **IT Help Desk & Manual Operator Agent:** Indexes local technical documentation, tool schemas, and hardware specs to debug operational failures and suggest system repairs.
- **Web Ingestion & News Monitor Agent:** Tracks breaking updates in local open-weight models, new tools, and upstream repository changes.
- **NPU / Local Inference Manager:** Balances weight offloading between Snapdragon Hexagon NPU, Jetson CUDA cores, and fallback cloud APIs.

6. GCP Cross-Account IAM Handshake & RLVR Loop

- **Business / Contractor Resource Tier:** Houses proprietary data in `gs://business-secure-vault-bucket/` and grants read-only roles/storage.objectViewer to the personal consumer service account.
- **Personal Developer Tier (\$1,000 Credit Pool):** Vertex AI Discovery Engine (`service-PERSONAL_PROJECT_NUMBER@gcp-sa-discoveryengine.iam.gserviceaccount.com`) points directly to the business bucket without copying data, absorbing 100% of RAG query costs.

- **The RLVR Training Loop:** Home Assistant packages verified task triples, Red Agent approves the batch, and the dataset streams to GCS. GCP instances run scheduled GRPO fine-tuning jobs on fixed base weights (Qwen/Llama) using Unsloth/Axolotl, returning fine-tuned GGUF weights back down to Node Alpha and Beta.

7. Decentralized Multi-Repository Map

master_workspace/	
├─ horizons-ui-v2.0/	# Kotlin/Java APK 1: Orchestration GUI, Webview, Chat Tile[cite: 8, 9]
├─ nova-daemon-shell/	# Kotlin/Java APK 2: OS Assistant & Accessibility Sandbox Bypass[cite: 8,
├─ nova-daemon-media/	# Kotlin/Java APK 3: Screen Vision, Moonshine STT, Kokoro TTS[cite: 1, 8,
├─ file-management-and-skills/	# Master Canonical Docs, JSONL Markers, Tool Schemas[cite: 1, 4]
├─ agent-harness-hub/	# Prime Agent (RLM), Pocock Skills, Honey for Devs, ECC, GSD[cite: 1, 7]
├─ localai-omni-route/	# Modality Routing, OpenAI/Anthropic Translation Endpoints[cite: 1, 6]
├─ obsidian-knowledge-vault/	# Graphify Networks, KAG/RAG Markdown Trees, NotebookLM Hooks[cite: 1]
├─ red-agent-auditor/	# Air-gapped Batch Training Verification, Nope Data Bank[cite: 1]
├─ node-beta-jetson/	# Postgres OB1 Vector Store, CUDA Inference Servers[cite: 1]
└─ gcp-training-flywheel/	# RLVR Export Scripts, Cross-Account IAM Handshake, GRPO Pipelines[cite: .

8. Universal JSONL Marker Schema

Every markdown documentation file and skill repository includes an attached .jsonl schema block to enable rapid, low-token search indexing across local agents and Vertex AI Discovery engines:

```
{
  "record_id": "SKILL_POCKOCK_GRILL_001",
  "document_path": "agent-harness-hub/pocock-skills/grill-me.md",
  "category": "SKILL_HARNESS",
  "target_runtime": "PRIME_AGENT_RLM",
  "metadata": {
    "title": "Grill-Me Architectural Interview Harness",
    "description": "Interactive question-driven prompt for technical planning",
    "primary_tools": ["claude-code", "open-wiki-cli", "prime-agent"],
    "required_context_keys": ["repo_scope", "target_architecture"]
  },
  "retrieval_tokens": ["grill", "spec", "architecture", "audit", "pocock"],
  "entry_points": {
    "repl_command": "/skill run grill-me",
    "jsonrpc_method": "agent.skills.execute"
  }
}
```

User prompt: Oh yeah, and Groddox is going to be priority one, bro. That's what's going to actually like make the plan for exactly which skills go where, what agents are going to be equipped with what, how we actually divide the files up and parse those out and what gets kept and you know, what's on the hot plate and what's in the vault.

Response: **AESOP XI: Autonomous Edge-Computing Architecture Blueprint**

Heterogeneous System Orchestration, Asymmetric Routing, & Autonomous Agent Flywheels

Specification Version: v2.2-Production-Ready | Target: Multi-Node Sovereign Mesh

1. Canonical Build Sequence & Execution Roadmap

- Phase 1: Grill with Docs Master Architectural Planning & Repo Partitioning
- Phase 2: Global Knowledge Curation, File Parsing, & Universal JSONL Marker Indexing
- Phase 3: Tool Harness Ingestion (Pocock, Honey for Devs, Prime Agent, ECC, GSD)
- Phase 4: 3-APK Engine Deployment (Horizons UI + Shell Daemon + Media Daemon)
- Phase 5: Local Dual-Model Inference & Continuous Improvement RLVR Loop (GCP)
- Phase 6: Multi-Node Mesh Orchestration & Autonomous Agent Swarm (Jetson / Pi)


- **Priority 1: "Grill with Docs" Master Architectural Session & Strategic Roadmap**
 - **The Orchestration Driver:** Execute the master "Grill with Docs" session first to interrogate baseline documentation and establish the complete technical battle plan.
 - **Repository & Directory Partitioning:** Formulate the precise directory schemas, define strict repository isolation barriers, and assign individual agent manifests.


- **Skill-to-Agent Allocation:** Map exactly which skills, harnesses, and tool endpoints are assigned to which agent runtimes (Node Alpha, Beta, Gamma, or Delta).
- **Data Tiering Strategy:** Define explicit boundaries for what gets parsed, what stays on the active "hot plate" (mem0 context cache and local active reasoning ledgers), and what gets archived into the deep cold vault (Obsidian markdown vaults and Postgres OB1 vector stores).
- **Priority 2: Global Knowledge Curation & Universal JSONL Marker Indexing**
 - Assemble, clean, and partition core technical documentation: Qualcomm QAIRT/HTP SDK, Android Media/Accessibility APIs, Google Vertex AI Agent Builder, Llama Server, and Unsloth.
 - Attach dense JSONL marker schemes to every markdown file, tool repository, and skill guideline to allow direct KAG/RAG indexing without context bloat.
- **Priority 3: Tool Harness Ingestion & Action Schemas**
 - Transform static reference repositories into machine-executable skills.
 - Ingest developer harnesses: **Matt Pocock Skills** (Grill-Me, Spec, Ticket, TDD, Code-Review), **Honey for Devs**, **Prime Agent** (RLM persistent REPL + Continual Harness with rollback snapshots), **ECC**, **GSD (Get Shit Done)**, **Claude Video**, **Reverse Skills**, **Code Review Graph**, **mem0**, and **OpenAI/Anthropic LocalAI Omni-Route bridges**.
- **Priority 4: 3-APK Native Android Deployment (Parallel with Priority 5)**
 - Deploy the decoupled 3-APK Android execution suite: **Horizons UI** (Orchestrator), **Shell Daemon APK** (Termux-free OS assistant/accessibility access), and **Media Daemon APK** (Screen-vision + Silero VAD + Speech stack).
- **Priority 5: Dual-Model On-Device Inference & Concierge Gateway**
 - Deploy the dual-model tandem: **Qwen 3.5 0.8B GGUF Q4_0** (Executor) and **Qwen 3.5 9B GGUF Q4_0** (Query / Metaprompt compiler).
 - Wire the on-device **Concierge Flow**: Voice capture → Silero VAD → Screen Vision context → Synthesized clean markdown meta-prompt → User verification → Frontier Model / LocalAI → Kokoro TTS speech output.
- **Priority 6: Network-Wide Integration & GCP Cloud Training Flywheel**
 - Establish the continuous improvement training flywheel: Real-time inference triples aggregated locally by Home Assistant → Packaged for batch audit → Validated by Red Agent → Exported to GCS → GCP Vertex AI / Compute Instances run RLVR (GRPO/Unsloth) fine-tuning → Updated weights pushed back to edge.
 - Activate the multi-agent home node: **Home Assistant Auditor & Housekeeper**, **Red Agent Batch Training Auditor**, **IT Help Desk & Manual Operator**, **Web Ingestion & News Monitor**, and **NPU Inference Manager**.

2. Multi-Node Compute Infrastructure

Node ID	Physical Platform	Primary Function	Core Acceleration & Compute	Primary Software & Daemons
NODE ALPHA	Moto Razr Ultra 2025	Local Orchestration & Concierge Hub	Snapdragon 8 Elite (Hexagon NPU v79)	3-APK Unified Stack (Horizons + Daemons), LocalAI, Moonshine/Kokoro
NODE BETA	Nvidia Jetson Orin Nano Super	Headless Home Server & Vector Hub	Dedicated CUDA Core Array (60–70 TOPs)	Postgres OB1 Protocol, Home Assistant, Red Agent Auditor, Multi-Agent Swarm
NODE GAMMA	Rubik Pi 3 (Dragonwing)	Display Workstation & Helpdesk Gateway	Thundercomm / Qualcomm SoC (14+ TOPs)	Dual-Monitor Display Engine, System Log Viewer, IT Manual Search
NODE DELTA	Google Cloud Platform	Asymmetric RAG & Heavy RLVR Training	Cloud TPUs / A100/H100 Instances	Vertex AI Agent Builder (\$1,000 Credit Tier), GCS Buckets, Unsloth/Axolotl GRPO

3. The 3-APK Unified Android Architecture





4. End-to-End System Dataflow & Runtime Routing



▼
[GCP TRAINING FLYWHEEL]
- Vertex AI Discovery Engine
- Unsloth/Axolotl GRPO Jobs
- New Weights Pushed to Edge

5. Multi-Agent Home Node Topology (Node Beta / Jetson)

- **Home Assistant (Cross-Agent Auditor & Housekeeper):** Manages real-time cross-agent auditing during live inference sessions, aggregates daily execution logs, compiles draft automation scripts, and packages raw telemetry pairs for downstream training.
- **Red Agent Auditor (Batch Training Guardrail):** Operates strictly when accumulated context reaches the training compilation threshold, conducting air-gapped safety verification against the Nope Data Bank before datasets are shipped to GCP.
- **IT Help Desk & Manual Operator Agent:** Indexes local technical documentation, tool schemas, and hardware specs to debug operational failures and suggest system repairs.
- **Web Ingestion & News Monitor Agent:** Tracks breaking updates in local open-weight models, new tools, and upstream repository changes.
- **NPU / Local Inference Manager:** Balances weight offloading between Snapdragon Hexagon NPU, Jetson CUDA cores, and fallback cloud APIs.

6. GCP Cross-Account IAM Handshake & RLVR Loop

- **Business / Contractor Resource Tier:** Houses proprietary data in `gs://business-secure-vault-bucket/` and grants read-only roles/storage.objectViewer to the personal consumer service account.
- **Personal Developer Tier (\$1,000 Credit Pool):** Vertex AI Discovery Engine (`service-PERSONAL_PROJECT_NUMBER@gcp-sa-discoveryengine.iam.gserviceaccount.com`) points directly to the business bucket without copying data, absorbing 100% of RAG query costs.
- **The RLVR Training Loop:** Home Assistant packages verified task triples, Red Agent approves the batch, and the dataset streams to GCS. GCP instances run scheduled GRPO fine-tuning jobs on fixed base weights (Qwen/Llama) using Unsloth/Axolotl, returning fine-tuned GGUF weights back down to Node Alpha and Beta.

7. Decentralized Multi-Repository Map

```
📁 master_workspace/
├── 📁 file-management-and-skills/ # Master Canonical Docs, JSONL Markers, Tool Schemas (Priority 1 Core)[cite: 8, 9]
├── 📁 horizons-ui-v2.0/           # Kotlin/Java APK 1: Orchestration GUI, Webview, Chat Tile[cite: 8, 9]
├── 📁 nova-daemon-shell/         # Kotlin/Java APK 2: OS Assistant & Accessibility Sandbox Bypass[cite: 8, 9]
├── 📁 nova-daemon-media/         # Kotlin/Java APK 3: Screen Vision, Moonshine STT, Kokoro TTS[cite: 1, 8, 9]
├── 📁 agent-harness-hub/         # Prime Agent (RLM), Pocock Skills, Honey for Devs, ECC, GSD[cite: 1, 7]
├── 📁 localai-omni-route/        # Modality Routing, OpenAI/Anthropic Translation Endpoints[cite: 1, 6]
├── 📁 obsidian-knowledge-vault/   # Graphify Networks, KAG/RAG Markdown Trees, NotebookLM Hooks[cite: 1]
├── 📁 red-agent-auditor/         # Air-gapped Batch Training Verification, Nope Data Bank[cite: 1]
├── 📁 node-beta-jetson/         # Postgres OB1 Vector Store, CUDA Inference Servers[cite: 1]
└── 📁 gcp-training-flywheel/     # RLVR Export Scripts, Cross-Account IAM Handshake, GRPO Pipelines[cite: 1]
```

8. Universal JSONL Marker Schema

Every markdown documentation file and skill repository includes an attached `.jsonl` schema block to enable rapid, low-token search indexing across local agents and Vertex AI Discovery engines:

```
{
  "record_id": "SKILL_POCKOCK_GRILL_001",
  "document_path": "agent-harness-hub/pocock-skills/grill-me.md",
  "category": "SKILL_HARNESS",
  "target_runtime": "PRIME_AGENT_RLM",
  "metadata": {
    "title": "Grill-Me Architectural Interview Harness",
    "description": "Interactive question-driven prompt for technical planning",
    "primary_tools": ["claude-code", "open-wiki-cli", "prime-agent"],
    "required_context_keys": ["repo_scope", "target_architecture"]
  },
  "retrieval_tokens": ["grill", "spec", "architecture", "audit", "pocock"],
  "entry_points": {
    "repl_command": "/skill run grill-me",
    "jsonrpc_method": "agent.skills.execute"
  }
}
```