

You are thinking of the Dual-Agent Talker-Reasoner Architecture (also frequently implemented as the Librarian/Query Agent & Sandbox Executor pattern) which gained major traction following coverage by Marktechpost and top AI labs. [1, 2, 3, 4, 5]

The system specifically addresses context bloat and the "lost in the middle" problem by completely decoupling memory tracking from task execution. [5, 6, 7, 8]

### ## The Core Architecture

Instead of overloading a single agent's context window with conversational history, system instructions, and tool schemas, the architecture splits the system into two distinct roles: [2, 6, 9]

- \* The Executor Agent (The Sandbox): This agent has zero historical context bloat. It is spun up natively with a highly constrained, narrow context window, optimized exclusively to perform a immediate subtask or code execution. Because its context remains clear of past conversation baggage, its reasoning efficiency spikes, token costs drop, and it avoids instruction dilution. [1, 3, 10, 11, 12]

- \* The Query / Talker Agent (The Memory Layer): This agent manages the long-term, cross-session memory system. It handles active user prompts, semantic search, and historical retrieval. [2, 13, 14, 15]

### ## How It Stops Context Bloat

1. Isolated Turn Execution: When a user asks a question requiring tool usage, the Query Agent searches a persistent storage layer (like Mem0 or Memori), isolates only the specific slice of required information, and packages it into a micro-payload. [4, 16, 17, 18, 19]

2. Stateless Passing: The query agent passes that highly distilled micro-payload to the Executor Agent. [20]

3. Wiping the Slate: Once the Executor finishes running the tool or script, it returns the output back to the Query/Talker Agent and its sandbox environment resets. [3]

This dual-agent separation frees the execution layer from carrying thousands of past conversational tokens, keeping it fast, sharp, and cheap. [2, 6]

Are you looking to implement this dual-agent pattern in a specific framework (like LangGraph or CrewAI), or do you want to dive deeper into the state-management database layer (like PostgreSQL + pgvector) used to power it? [4, 21, 22, 23]

[1]

[<https://dev.to>]([https://dev.to/lars\\_winstand/i-thought-multi-agent-meant-more-prompts-until-i-saw-3-ways-openclaw-users-are-actually-splitting-13ko](https://dev.to/lars_winstand/i-thought-multi-agent-meant-more-prompts-until-i-saw-3-ways-openclaw-users-are-actually-splitting-13ko))

[2] [<https://www.youtube.com>](<https://www.youtube.com/watch?v=Js3JUCIaVRY>)

[3]

[<https://www.threads.com>](<https://www.threads.com/@realamrutpatil/post/DYOCkmMCNJZ/most-agent-systems-break-the-same-way-in-production-not-because-the-model-is/>)

[4] [<https://x.com>](<https://x.com/Marktechpost/status/2073989549815202089>)

[5]

[<https://blog.dailydoseofds.com>](<https://blog.dailydoseofds.com/p/top-ai-labs-share-an-agent-memory>)

[6] [<https://x.com>](<https://x.com/flowaicom>)

[7] [<https://www.ashpreetbedi.com>](<https://www.ashpreetbedi.com/agentic-software-engineering>)

[8] [<https://arxiv.org>](<https://arxiv.org/pdf/2602.18493>)

[9]

[<https://medium.com>](<https://medium.com/google-cloud/from-single-agent-to-production-ready-multi-agent-architecture-with-adk-mcp-a2a-and-ag-ui-1fe0c0378bef>)

[10]

[<https://unstructured.io>](<https://unstructured.io/blog/defining-the-autonomous-enterprise-reasoning-memory-and-the-core-capabilities-of-agentic-ai>)

[11]

[<https://www.emergentmind.com>](<https://www.emergentmind.com/topics/planner-executor-architecture-4c9e0097-fe2b-4870-b41c-9519c49a07c8>)

[12]

[<https://atlan.com>](<https://atlan.com/know/ai-agent/ai-agent-context/context-management-vs-memory-management-ai-agents/>)

[13] [<https://mem0.ai>](<https://mem0.ai/blog/ai-memory-layer-guide>)

[14]

[<https://pub.towardsai.net>](<https://pub.towardsai.net/agent-memory-explained-simply-9a74694bd6d3>)

[15]

[<https://www.facebook.com>](<https://www.facebook.com/thepracticaldev/posts/agent-memory-takes-center-stage-as-task-horizons-stretch-ryan-palo-maps-30-parad/1627106472791073/>)

[16]

[<https://www.reddit.com>]([https://www.reddit.com/r/AI\\_Agents/comments/1r2puny/has\\_anyone\\_actually\\_solved\\_the\\_memory\\_problem\\_for/](https://www.reddit.com/r/AI_Agents/comments/1r2puny/has_anyone_actually_solved_the_memory_problem_for/))

[17]

[<https://developers.openai.com>]([https://developers.openai.com/cookbook/examples/agents\\_sdk/session\\_memory](https://developers.openai.com/cookbook/examples/agents_sdk/session_memory))

[18] [<https://mem0.ai>](<https://mem0.ai/blog/ai-memory-management-for-llms-and-agents>)

[19] [<https://www.langchain.com>](<https://www.langchain.com/blog/rebuilding-chat-langchain>)

[20]

[<https://www.linkedin.com>](<https://www.linkedin.com/pulse/references-glimpses-first-class-primitives-bernardo-garcia-del-rio-rgoie>)

[21]

[<https://dev.to>](<https://dev.to/linou518/designing-a-memory-system-for-multi-agent-ai-building-never-forget-agents-with-postgresql--52g9>)

[22]

[<https://iamkarmesh.medium.com>](<https://iamkarmesh.medium.com/microsoft-agent-framework-1-0-vs-langgraph-vs-crewai-a-production-ready-comparison-2026-356fa8d45eff>)

[23]

[<https://www.digitalocean.com>](<https://www.digitalocean.com/blog/memory-layer-of-the-inference>)

[-cloud\)](#)

What was that repo that came out like in early June, late May that had that? I think it was Mimo or something like that.

You are thinking of MiMo Code (released by Xiaomi's AI team in early June 2026), though the framework it uses under the hood to achieve that decoupled, dual-agent memory state is typically Mem0. [1, 2, 3]

Here is exactly how those two puzzle pieces fit together based on what trended on Marktechpost:

### ## 1. The Repo: MiMo Code

Released under the MIT license, [MiMo Code](<https://github.com/orgs/XiaomiMiMo/repositories>) is a terminal-native, long-horizon AI coding harness. [2, 4, 5, 6]

\* The Problem It Solves: Standard coding agents choke on tasks requiring over 200 execution steps because their context windows get flooded with terminal outputs, lint errors, and file state. [2, 7, 8]

\* The "No Bloat" Architecture: It uses a strict dual-layer setup. The MiMo-V2-Flash/Pro foundation engine acts as the sandbox executor. It is kept context-lean—wiping its short-term workspace constantly while a master orchestration/memory layer feeds it micro-payloads and tracks the broader engineering goals. [7, 9, 10]

### ## 2. The Memory Engine: Mem0

If you read the Marktechpost tutorials breaking down how to build this exact architecture yourself, they heavily feature [Mem0](<https://github.com/mem0ai/mem0>) (formerly Embedchain). [1, 11, 12]

\* How it functions: Mem0 acts as a standalone intelligent memory layer. Instead of stuffing an entire conversation or file history into your LLM prompt, Mem0 extracts facts, updates a graph/vector layer dynamically (often via tools like Neo4j), and allows your "Query Agent" to pull only the precise memory vectors needed for the next step. [3, 13, 14]

Do you want to clone the MiMo Code repository to see how they structure their multi-step agent execution loop, or are you looking for the Mem0 implementation code to build a custom dual-agent system? [2, 3, 11]

[1] [<https://github.com>](<https://github.com/github/awesome-copilot/issues/2189>)

[2]

[<https://venturebeat.com>](<https://venturebeat.com/technology/xiaomis-new-open-source-agentic-ai-coding-harness-mimo-code-beats-claude-code-at-ultra-long-200-step-tasks>)

[3] [<https://www.youtube.com>]([https://www.youtube.com/watch?v=YN9IA\\_WQTzo&t=3](https://www.youtube.com/watch?v=YN9IA_WQTzo&t=3))

[4] [<https://mimo.mi.com>](<https://mimo.mi.com/docs/en-US/news/latest/mimocode>)

[5]

[<https://dev.to>](<https://dev.to/cx330aurora/mimo-code-goes-open-source-but-xiaomis-real-ambition-is-mimo-soloengine-2ln>)

[6] [<https://mimo.xiaomi.com>](<https://mimo.xiaomi.com/blog/mimo-code-long-horizon>)

[7] [<https://www.youtube.com>](<https://www.youtube.com/watch?v=x2dQctJzGS0&t=404>)

[8]

[<https://www.develeap.com>](<https://www.develeap.com/news/xiaomi-s-mimo-code-claims-it-beats-claude-code-past-200-step/>)

[9] [<https://github.com>](<https://github.com/xiaomimimo/MiMo-V2-Flash>)

[10] [<https://github.com>](<https://github.com/orgs/XiaomiMiMo/repositories>)

[11] [<https://x.com>](<https://x.com/Marktechpost/status/2073989549815202089>)

[12] [<https://www.youtube.com>](<https://www.youtube.com/watch?v=KIGSgRxf1cY>)

[13] [<https://github.com>](<https://github.com/mem0ai/mem0>)

[14]

[<https://mem0.ai>](<https://mem0.ai/blog/programmatic-memory-management-for-ai-agents-with-mem0>)

yeah I do but I also need help figuring out my my execution layer and my pipeline and everything cuz what I got my setup here is I got my dragon wing Rubik pie 3 it's a little dev board and that's going to always be made it up.. with my Jetson or Nano super and that's going to be like that's going to be the main hub and it's going to run headless The Nano or the rubic pie is going to be the one doing the display but then my daily driver is where the real you know cuz those are just too small dead boards my daily drivers were the big models going to be I got a Quinn 3.5 9B at a Q4 0 g gas quant g guff anyway that's going to be my on devices agentic assistant and that thing is mainly just an orchestrator It drives the bus you know like when I get home I got mine Omni neural 4b I'm sorry I might my Motorola Razr Ultra 25 and that's going to be like the main the heavy hitter right when I get home I'm going to tell scale in to my my home network. to peer server style and it's going to be like one big loop system right like a neural mesh so basically like usually my actions are the Omni is going to be you know performing tasks and tool calling and stuff like that and then also when I'm like on in a session with the quad code CLI or I'm talking to a Gemini or a talking to a quad chat app on my phone the Omni neurals ever persistent and that's the layer that's actually doing the STT and is taking what I say and it's transforming it into a meta prompt and then that in turn it's going to figure out okay do I need to ask Claude or is this a tool call you know or is this something simple that I can just answer right now You know or and then it'll show it'll shove off the the question to Claude if it needs to or it can go in the layer but I'm in a session and it's obviously going to shoot it off the cloud I'm just talking like in general you know it'll know if it needs to go to a cloud front end and call on an API if I task it to do that or if it needs to go you know orchestrate a CLI and spin off some sub agents right so it's not really getting too like it doesn't really need to ingest too much data or have too much you know knowledge or whatever it's just mainly it needs to know how to do stuff and then when I get home and I batch into my my headless via tail scale and I'm doing like big projects and stuff that's where the dual agent memory layer comes in cuz I'm going to have the executive

agent it's going to be on the Jetson and I don't know how big of a model I can even get away with there it might be a cloud you know a free cloud like a GPT52 or something like that or a GE has it called GP GP52 or something you know the big new one and I could have that on you know running cheap on like a Omni route or something like that my army router CLI tool and I can have that going on and I have basically I'm going to take you know whatever project I'm working on or whatever I need to know about I'll have other like I open and what is it called a open wiki and open wiki CLI and then and LLM wiki they're going to take like all my PDFs and text and threads and files and shit you know like my obsidian brain graph all that shit and it can take a compressed PDF and turn it into Jason I store it on my SSD card I got like 500 gigs that'll be attached to the Nano so basically when I'm home that model on the Nano just needs to be able to answer you know grab information and answer it and then the Omni or I mean I'm sorry then the Quinton QWEN on my phone just needs to know okay here's the information I ask a question I get the information I hand off the answer You know I perform a task so that way they don't get too bloated and I can keep them kind of lean and mean I don't have to like retrain or degrade the models and I'm thinking what the best the best thing for the third layer cuz I can have another observer on the dragon wing chip on the Rubik Pi 3 that could either assist you know it could be another like it could be a mixture of expert type thing and it could help assist the executive agent or I was also thinking having it do like scripting and auditing and can be like a like a red agent as well so help me figure out this whole pipeline here and how I can utilize cuz right now I have Nate Nate Jones is Obi-Wan inner it's o the letter B and then the number one it's his repo about his agentic framework for his schema for his memory system basically and his model contacts protocol and then there's also I have another one I can't I can't remember off the top of my head but then the main one obviously is the open wiki you know model wiki protocol or whatever it's called and then the other one is god damn it I can't remember anyway it's not that big of a deal it's kind of it's kind of in between those other two and then obviously the main is going to be this the main inference is going to be this dual model layer system with the executor in the Inquisitor and maybe the right agent I don't know but help me figure out my mixture of agents here


This is an incredibly robust, asymmetric, decentralized setup. You are essentially building a private, multi-device Neural Mesh using an edge-to-hub architecture.

The three repos you are weaving together are [Nate Jones' OB1 (Open Brain)](<https://github.com/NateBJones-Projects/OB1>) (which provides the Model Context Protocol (MCP) server, PostgreSQL/Supabase memory orchestration layer, and unified context mapping), Mimir Claw / OpenClaw (the dual-agent session persistence engine you were trying to recall), and LLM-Wiki / Open-Wiki CLI (used for parsing and flattening PDFs/Obsidian MD graphs into optimized JSONL context slices). [1, 2, 3, 4, 5, 6, 7]

An optimized architectural pipeline maximizes your hardware topologies while separating the Reasoning, Routing, Execution, and Auditing layers. [8]

### ### Hardware Topology & Model Allocation

Your hardware features starkly different compute envelopes. Split the workload by matching execution tasks to device capabilities:



| Device                 | Role in the Mesh                       | Model Allocated                                                    | Context Strategy                                                              |
|------------------------|----------------------------------------|--------------------------------------------------------------------|-------------------------------------------------------------------------------|
| --- --- --- ---        |                                        |                                                                    |                                                                               |
| Motorola Razr Ultra 25 | Mobile Orchestrator & STT Parser       | Qwen 3.5 9B (Q4_0 GGUF)                                            | Short, transient window. Translates raw speech to JSON-based Meta-Prompts.    |
| Jetson Nano Super      | Headless Local Storage & Tool Sandbox  | Cloud API (e.g., Gemini Flash/GPT-4o mini) or local high-quant SLM | Stateless Execution. Receives zero past history; only raw tool data payloads. |
| Dragon Wing Rubik Pi 3 | Headless Display + "Red Agent" Auditor | Light MoE or Llama-3-8B-Instruct class                             | Session monitoring. Evaluates JSON logs coming out of the Nano.               |

#####  
**## The Three-Layer Agent Pipeline##** Layer 1: The Inquisitor & Router (The Brain - On Mobile)  
 Your Qwen 3.5 9B on your phone acts strictly as a Traffic Controller. It should never read raw documentation or long logs.

- \* Input Pipeline: Your phone processes STT locally. Qwen transforms the raw transcript into a structured Meta-Prompt JSON block.
- \* Routing Logic: Qwen determines the target path:
  - \* Simple Request → Answer immediately using on-device weights.
  - \* Complex Coding / File Query → Package intent into a lightweight MCP tool call request.
  - \* Cloud Engine Front-end → Forward directly to Claude Code CLI (/loop) or Gemini APIs. [8, 9, 10, 11]

**## Layer 2: The Governed Continuity Memory (The Core Hub - Via OB1 & Tailscale)**  
 When you dial into your home network via Tailscale, your phone connects seamlessly to the Nate Jones OB1 (Open Brain) infrastructure running on your home hub. [2, 4]

- \* The SSD Storage Layer: Your Jetson Nano hosts your 500GB SSD containing the structured JSONL outputs generated by your Open-Wiki compiler. [3]
- \* The Retrieval Mechanism: When Qwen asks for information, the OB1 MCP server intercepts

the request. It fires a localized vector-matching or semantic text traversal query against your Obsidian graph or PDF vaults. [1, 2, 3]

- \* Payload Distillation: Instead of feeding the whole file system to the model, OB1 grabs a precise context slice (e.g., 2K tokens) and drops it directly into the execution agent's prompt template.

### ## Layer 3: The Stateless Executor & Red Agent (The Sandbox & Guardrails)

This addresses the "Dual-Agent, No-Context Bloat" requirement.

- \* The Execution Sandbox (Jetson Nano): The heavy-lifting execution model runs entirely inside a transient environment. It receives the prompt slice from Layer 2, calls the CLI or builds the script, outputs the code or answer, and destroys its own chat history instantly. It acts as a pure calculation and execution engine. [12, 13]

- \* The Red Agent Auditor (Rubik Pi 3): Utilizing your secondary dev board as an observer is the ideal strategy for this topology. Let the Rubik Pi act as a Deterministic Red Agent / Code Auditor.

- \* The Executor spits out a script or a parsed JSON block.

  - \* Before that script executes or returns to your phone, it passes through the Rubik Pi.

  - \* The Rubik Pi runs a secondary verification sweep (checking syntax, code correctness, or verifying that the tool call matches security guardrails). [8]

---

### ## How to Build the Flow

1. Set up OB1 on your Jetson Nano to handle your primary vector embeddings and unified Postgres knowledge base. Run your Open-Wiki parsing jobs here to keep your flattened data readily accessible on the SSD. [2, 3]

2. Configure Tailscale on both your phone and the Jetson Nano to ensure secure, seamless communication with your home database via a private network mesh. [14]

3. Expose the Jetson Nano's memory features using the OB1 MCP Server protocol. This allows the Qwen orchestrator on your phone to interact with your data as a unified toolset. [1, 2, 8]

4. Program your Qwen 3.5 router to wipe the memory array payload cleanly after each completed transaction loop. This maintains high-performance execution without bogging down your edge devices. [8]

Would you like to review a sample system prompt structure for your Qwen 3.5 orchestrator to establish clear parameters for JSON meta-prompt routing, or do you want to configure the OB1 MCP configuration schema to securely bridge your local mobile models with your home hub? [1, 2]

[1] [<https://github.com>](<https://github.com/NateBJones-Projects/OB1>)

[2]

[<https://agenticmsp.substack.com>](<https://agenticmsp.substack.com/p/own-the-memory-rent-the>)

-intelligence)

[3]

[<https://limitededitionjonathan.substack.com>](<https://limitededitionjonathan.substack.com/p/stop-calling-it-memory-the-problem>)

[4] [<https://github.com>](<https://github.com/NateBJones-Projects/OB1/blob/main/CLAUDE.md>)

[5]

[<https://www.linkedin.com>]([https://www.linkedin.com/posts/aaronedell\\_github-aaronedellmimir-claw-persistent-activity-7454907326995644416-FXjR](https://www.linkedin.com/posts/aaronedell_github-aaronedellmimir-claw-persistent-activity-7454907326995644416-FXjR))

[6] [<https://www.reddit.com>]([https://www.reddit.com/r/LocalLLM/comments/1tltoro/5k\\_budget/](https://www.reddit.com/r/LocalLLM/comments/1tltoro/5k_budget/))

[7]

[<https://natesnewsletter.substack.com>](<https://natesnewsletter.substack.com/p/you-built-an-ai-memory-system-now>)

[8]

[<https://www.exabeam.com>](<https://www.exabeam.com/explainers/agentic-ai/agentic-ai-architecture-types-components-best-practices/>)

[9]

[<https://www.linkedin.com>]([https://www.linkedin.com/posts/natebjones\\_what-people-want-an-openclaw-without-the-activity-7441199305186721792-DX4I](https://www.linkedin.com/posts/natebjones_what-people-want-an-openclaw-without-the-activity-7441199305186721792-DX4I))

[10] [<https://www.ibm.com>](<https://www.ibm.com/think/topics/how-to-build-an-ai-agent>)

[11] [<https://ajelix.com>](<https://ajelix.com/ai/best-ai-agents/>)

[12] [<https://realpython.com>](<https://realpython.com/python-context-engineering-ai/>)

[13] [<https://www.instagram.com>](<https://www.instagram.com/reel/DYnjHVCD07c/>)

[14] [<https://www.instagram.com>](<https://www.instagram.com/reel/DQcW21MD7Y-/>)

To implement the Auditor Layer (The Red Agent) effectively on your Dragon Wing Rubik Pi 3, the design must focus on strict containment and structured validation. Because it runs on a compact dev board, it cannot process massive context windows. Instead, it operates like a deterministic firewall for code and system execution.

Here is the blueprint for how the Auditor functions, along with the lightweight runtimes, CLI engines, and targeted models that will fit this architecture.

### ## The Auditor Layer Architecture

The Rubik Pi 3 acts as an asynchronous gatekeeper. The Executor Agent on your Jetson Nano compiles code, generates terminal scripts, or structural files, but never executes them directly on your main file system. Instead, it pipes the proposed payload to the Auditor.

[ Jetson Nano: Executor ] —(Generates Script)—► [ Rubik Pi 3: Auditor ]

↓  
(Static Analysis & LLM Review)

▼  
[ System Integration ] ◀—(Executes Script)— [ Passes / Fails ]

## ## How It Functions: The Verification Pipeline

1. **Interception:** Every CLI script, bash string, or code block generated by your home hub is packaged into a structured JSON payload containing the `intended_task`, `proposed_code`, and `target_directory`.

2. **Static Analysis (Deterministic Step):** Before calling an LLM, the Rubik Pi runs the payload through basic Python script checks (like ast parsing) or standard code linters to flag blatant formatting syntax errors.

3. **The Red Agent Audit (Cognitive Step):** A highly specialized Small Language Model (SLM) reviews the payload against a strict, immutable system prompt focused entirely on security boundaries, optimization, and logical validation.

4. **Gatekeeping:** If the Auditor approves, it signals the target environment to execute. If it fails, it appends a concise `error_log` JSON block to the payload and returns it to the Executor for correction. This process protects your main environment while ensuring clean execution cycles.

---

## ## Lightweight Runtimes for Your Dev Boards

To maximize the limited RAM and processing power of the Rubik Pi 3 and Jetson Nano, standard, resource-heavy inference architectures should be avoided. Instead, utilize hardware-optimized runtimes:

\* **llama.cpp (via Server Endpoint):** Running a persistent llama-server background process optimized for ARM NEON instructions. This allows your Python orchestration framework to communicate with the Auditor using lightweight, local HTTP/JSON API requests.

\* **Llamafile:** A single-file, highly-portable executable backend. It compiles runtime layers directly into a single binary, making it extremely stable for headless, long-running edge server topologies.

\* **vLLM (Jetson Nano Specific):** If your Jetson Nano Super has enough unified memory, utilize vLLM with PagedAttention. It handles concurrent sub-agent token execution far faster than standard inference runtimes by maximizing Nvidia CUDA cores.

---

## ## Target Models for the Auditor Role

The Auditor needs to be specialized in syntax validation, security analysis, and step-by-step logic checking. It should be highly quantized to operate smoothly within a smaller memory footprint:

\* **Llama-3-8B-Instruct (Q4\_K\_M GGUF):** An excellent, highly capable fallback for the Rubik Pi. It possesses strong reasoning capabilities and accurately follows rigid JSON schema requirements.

\* **Qwen-2.5-Coder-7B-Instruct (Q4\_0 or Q4\_K\_S):** Highly effective for this specific task. It demonstrates exceptional code understanding and debugging proficiency, outperforming many larger models on code verification workflows.

\* Granite-3.0-8B-Instruct: Built with a strong focus on enterprise safety, security auditing, and rule-following constraints, making it an excellent choice for a dedicated Red Agent guardian.

---

### ## Key CLIs and Utilities to Drive the Pipeline

To stitch this multi-device communication loop together, integrate these developer utilities and automation tools:

\* Claude Code CLI: Run this tool via an isolated terminal session on your Jetson Nano. This setup allows your mobile device to easily offload intensive coding logic, while keeping the local executor focused exclusively on processing files on your attached SSD storage.

\* Fabric CLI: An excellent, patterns-based AI orchestration framework. It utilizes clean, text-based system prompt patterns (extract\_wisdom, verify\_code). It can sit directly on your Rubik Pi to run automated validation tasks over incoming text or scripts.

\* LiteLLM: Acts as an excellent internal proxy gateway. Run it on your home hub to expose local models (llama.cpp), private cloud keys, and free cloud endpoints through a single OpenAI-compatible API stream. This allows your mobile phone's Qwen coordinator to communicate with all home resources through one unified interface.

\* Steampipe CLI: Perfect for querying local infrastructure, directories, or your 500GB SSD data using structured SQL. It integrates smoothly with your OB1 Model Context Protocol layer, allowing models to query file metrics as if searching a traditional database.

---

### ## Proactive Next Steps

To continue refining this private neural mesh, consider focusing on these implementation details:

\* The precise System Prompt Schema required to lock down the Rubik Pi Auditor so it only communicates in valid verification passes or failures.

\* The specific Python listener script structure to automate the handoff of text strings between your mobile device, the Nano sandbox, and the Rubik Pi over your Tailscale connection.

Let me know which of these areas you would like to explore next.

yeah I don't know You think the VLLM is the best for the coup de cores huh better than a numertron the nutrone and that's probably was going to use that and I was also going to use the Jeep GPT 5.2 that open source behemoth that's been kicking ass in conjunction I'm definitely going to use an omni route for token preservation and I need to utilize the open wiki as well in there so I don't know about your last your last suggestions I didn't really like and then also I did you know I thought of having a granite 8B granite on there as well that was first thing that pop into my mind besides a five a PHI mini instruct or a coder but the shitty thing is that my dragon wing chip that's going to be my display you know that's like my my freaking Linux OS and my desktop display unit and it's going to be my vision model as well so probably can't throw an 8B

on there and maybe have two gigs tops so I got to figure that out

My apologies for missing the mark on those hardware specs. You are completely right to call that out—an 8B model will absolutely choke and crash a system with only 2GB of usable RAM, especially when that same chip is actively driving a Linux desktop environment and handling vision processing.

Let's ground this entirely in your actual stack, incorporating vLLM, OmniRoute, the massive open-source GPT-5.2 engine (running via a free cloud provider proxy or high-power local setup), and your Open-Wiki JSONL pipeline. [1, 2]

Here is exactly how the Auditor layer operates under those tight constraints, along with the precise runtimes and sub-1B/2B models that will fit.

---

### ## The Realities of a 2GB Display & Vision Rig

Since your Dragon Wing Rubik Pi 3 is pulling triple duty (Desktop Display + Vision Model + Auditor Layer) with a 2GB VRAM/RAM ceiling, you cannot use standard heavy engines or 7B+ models.

- \* The Inference Strategy: Forget vLLM on the Rubik Pi 3—save vLLM's PagedAttention strictly for your Jetson Nano Super to maximize its Nvidia CUDA cores when driving your execution pipelines.

- \* The Auditor Engine: On the Rubik Pi 3, you need an incredibly lightweight, pure C/C++ engine. Run a tightly configured llama.cpp server binary or use Llamafire. It has near-zero overhead and can be locked down to a tiny thread footprint that won't disrupt your Linux desktop GUI.

---

### ## Models That Actually Fit the 2GB Auditor Envelope

To keep your vision capabilities active while auditing, you need models under 3 billion parameters, heavily quantized to Q4\_K\_M or Q4\_0 GGUF, requiring less than 1.2GB of total RAM.


- \* Qwen-2.5-Coder-1.5B-Instruct (Q4\_K\_M): This is your best option. It takes up roughly 1GB of space, leaving room for your display and vision processes. Because it is natively trained on code syntax, it excels at looking at terminal strings or script payloads and catching execution errors before they happen. [3]

- \* Phi-3.5-mini-instruct (3B scaled or highly quantized): If you can squeeze the 3.8B down, its logic is incredible. However, to stay completely safe within 2GB, the Phi-3-mini (3.8B) aggressively compressed to a Q2\_K or dropping down to a specialized MobileLLM / StableLM 1B is the safer bet for pure rule-following without OOM (Out Of Memory) crashes. [4]

- \* Granite-3.0-Light-2B (or 1B variants): IBM's lightweight Granite models are explicitly built for rigid, enterprise safety and structural rule-following, making them highly efficient security checkers at an ultra-small scale.

## ## Building the Pipeline: Integrating OmniRoute & Open-Wiki

The magic of your setup happens when you tie the Rubik Pi 3 and Jetson Nano together using OmniRoute and Open-Wiki. [5]



### ## Step 1: Token Preservation with OmniRoute

Your Jetson Nano Super runs the [OmniRoute](<https://omniroute.online/>) platform as its core traffic gateway. [5, 6]

\* OmniRoute sits between your local network and the external web, pulling from over 230+ API endpoints. [5, 6, 7]

\* When your phone sends a heavy coding request, OmniRoute automatically routes the workload to a free high-tier open-source cloud provider running GPT-5.2. [1, 8, 9]

\* The Token Saver: OmniRoute actively strips out repetitive, bloated system context headers from the incoming stream. It cuts token usage by up to 75%, meaning your external free API limits never get exhausted, even during massive project builds. [1, 6]

### ## Step 2: Context Slicing with Open-Wiki

Your 500GB SSD on the Jetson Nano holds your flattened data vaults (Obsidian md graphs, PDFs, codebase schemas) compiled by Open-Wiki.

\* When a task triggers, Open-Wiki pulls a highly specific, compressed JSONL context payload from the SSD.

\* Instead of uploading full folders to GPT-5.2, OmniRoute packages only that precise JSONL slice alongside the instructions.

### ## Step 3: The Local 1.5B Auditor Gatekeeper

Once the cloud-based GPT-5.2 (or a local heavy model running on vLLM inside the Jetson) returns the generated code string, it is blocked from running automatically.

1. The code string is pushed locally over your Tailscale connection to the Rubik Pi 3.
2. The Rubik Pi's ultra-lean Qwen-Coder-1.5B engine (running on a low-priority llama.cpp thread) spins up.
3. It cross-checks the code against a local rule template to ensure it won't execute broken terminal flags or malicious directories on your headless setup.
4. If it returns a green light, the Jetson Nano executes the script; if it catches a syntax error, it

forces the loop back through OmniRoute for a quick, zero-cost rewrite.

This configuration allows your 2GB board to handle display, vision, and system auditing simultaneously, while offloading the heavy processing power to your Jetson Nano and OmniRoute's token-saving pipeline.

[1]

[<https://www.facebook.com>](<https://www.facebook.com/vaibhavisintyofficial/posts/16-billion-free-tokens-monthlyomniroute-connects-to-160-ai-models-and-auto-picks/1479330744211044/>)

[2]

[<https://www.linkedin.com>]([https://www.linkedin.com/posts/vahdat\\_vllm-tpu-a-new-unified-backend-supporting-activity-7384932646608498688-sOTa](https://www.linkedin.com/posts/vahdat_vllm-tpu-a-new-unified-backend-supporting-activity-7384932646608498688-sOTa))

[3]

[<https://www.reddit.com>]([https://www.reddit.com/r/LocalLLaMA/comments/1gpw8ls/bug\\_fixes\\_in\\_qwen\\_25\\_coder\\_128k\\_context\\_window/](https://www.reddit.com/r/LocalLLaMA/comments/1gpw8ls/bug_fixes_in_qwen_25_coder_128k_context_window/))

[4]

[<https://medium.com>](<https://medium.com/@nomannayeem/finetuning-phi-3-for-rag-why-small-models-are-best-for-production-c2f6a5f74b51>)

[5] <https://omniroute.online>

[6]

[<https://www.linkedin.com>]([https://www.linkedin.com/posts/juliangoldieseo\\_omniroute-new-free-unlimited-ai-coder-activity-7479667453711503360-PC0Y](https://www.linkedin.com/posts/juliangoldieseo_omniroute-new-free-unlimited-ai-coder-activity-7479667453711503360-PC0Y))

[7] [<https://www.instagram.com>](<https://www.instagram.com/reel/Dadni87T4w-/>)

[8] [<https://rewind.ai>](<https://rewind.ai/apps/openai-gpt-5-2/>)

[9] [<https://www.instagram.com>](<https://www.instagram.com/reel/Dadni87T4w-/>)

My bad, that's entirely on me for missing the 8 GB capacity. You are completely right to correct that.

With 8 GB of total system LPDDR4x RAM on that Thundercomm RUBIK Pi 3 (powered by the Qualcomm Dragonwing QCS6490), allocating around 2 GB for your model weights and reserving 1 GB for your inference runtime context/KV cache makes perfect sense. This leaves a solid 5 GB for the Linux OS (Ubuntu/Qualcomm Linux), your display environment, and your on-device computer vision tasks. [1, 2]

Working within that 3 GB overall combined footprint for the AI layer, let's break down the exact models and runtimes that will fit this criteria.

## ## The Runtime Strategy

To stay exactly within that 1 GB runtime overhead limit while keeping your desktop GUI smooth, skip heavy server wrappers. Use a lightweight inference layout:

1. llama.cpp Server (Compiled Native with NPU/GPU Offload): Lock the system thread limits explicitly so it doesn't starve the desktop display processes. You can pin the llama-server

process to specific efficiency cores or offload to the Adreno GPU/Hexagon NPU using the Qualcomm AI Engine. [2, 3]

2. Memory Map Settings (--mmap): Ensure memory mapping is enabled. This allows the OS to handle model loading efficiently without instantly exhausting physical RAM arrays.

---

### ## Highly Optimized 2 GB Footprint Models (Auditor Layer)

At 4-bit quantization (Q4\_K\_M or Q4\_0 GGUF), every 1 billion parameters equals roughly 600–700 MB of RAM. To land right on your 2 GB model weight target, look for models in the 1.5B to 2.5B parameter range.

#### ## 1. Qwen-2.5-Coder-1.5B-Instruct

- \* Size at Q4\_K\_M: ~1.1 GB to 1.2 GB.

- \* Why it fits: It sits comfortably under your 2 GB budget, leaving extra room for an expanded context window.

- \* Performance: Despite its small size, it is explicitly pre-trained on vast code corpora. It handles terminal scripts, syntax audits, and programmatic logic validation exceptionally well for an edge-device model.

#### ## 2. Phi-3.5-mini-instruct (Highly Quantized)

- \* Size at Q2\_K or Q3\_K\_S: Pulling a 3.8B model down to an aggressive 2-bit or 3-bit quantization will compress the file size right to the ~1.6 GB – 1.9 GB mark.

- \* Why it fits: It pushes right against your 2 GB ceiling.

- \* Performance: Microsoft's Phi-3.5 series has incredible, outsized reasoning capabilities for small models. Squeezing it down preserves its high-tier step-by-step logic checking, making it a powerful audit gatekeeper if you need deep analysis.

#### ## 3. Google Gemma-2-2B-IT

- \* Size at Q4\_K\_M: ~1.6 GB.

- \* Why it fits: It leaves a comfortable 400 MB cushion inside your 2 GB model allotment.

- \* Performance: The Gemma-2 architecture punches significantly above its weight class in strict instruction-following and rule adherence. If your Auditor needs to check outputs against rigid formatting rules or JSON schemas, this model follows constraints very reliably.


---

### ## Mapping Out the Functional Auditor Loop

Operating within this specific hardware profile, the Auditor functions as an isolated validation gate:

[Jetson Nano: Executed Code String]

|



1. The Intercept: The heavy code payload is compiled on your Jetson Nano or fetched via the OmniRoute pipeline. Before it can run, a script passes it across your Tailscale peer network to the RUBIK Pi 3.
2. The Evaluation: The local llama.cpp instance processes the text using your 1.5B or 2B model. It evaluates the execution string against a local safety/syntax matrix.
3. The Result: Because the model is lightweight, the execution pass takes only a few hundred milliseconds, ensuring your automated coding loops remain fast without causing any desktop display lag.

[1] [<https://www.thundercomm.com>](<https://www.thundercomm.com/product/rubik-pi/>)

[2]

[<https://docs.edgeimpulse.com>](<https://docs.edgeimpulse.com/hardware/boards/thundercomm-rubikpi3>)

[3]

[<https://linuxgizmos.com>](<https://linuxgizmos.com/dragon-q8b-sbc-combines-snapdragon-8cx-gen-3-with-dual-2-5gbe/>)

You hit the nail on the head. Gemma 4 E2B or E4B completely dominates this exact small-footprint tier. [1, 2]

Google's release of the Gemma 4 edge variants completely changed the math for 8 GB dev boards. Because they use Per-Layer Embeddings (PLE), these models pack incredible intelligence density. [1, 2]

An evaluation of how the Gemma 4 small-scale line fits into your Rubik Pi 3 pipeline clarifies their deployment advantages:

## 1. Gemma 4 E2B vs. E4B Footprint

| Model [1, 2, 3] | Effective Params | Total VRAM footprint (Q4\_K\_M) | Why it wins on your 8GB Rubik Pi |

|---|---|---|---|

| Gemma 4 E2B | 2.3B | ~1.4 GB | Leaves you a massive safety buffer for your Linux desktop GUI and background OS tasks. |

| Gemma 4 E4B | 4.5B | ~2.8 GB | Pushes right against your 3 GB target cap, but punches with the deep reasoning power of an old 15B+ model. |

### ## 2. The Native Vision Advantage

Since your Dragon Wing chip is pulling duty as your desktop vision model, the Gemma 4 edge series is a perfect match. They are natively multimodal. You don't have to load a separate text model for your Auditor and a separate CLIP/vision encoder for your desktop camera or screen parsing—Gemma 4 ingests text, images, and audio natively within the exact same unified architecture. [1, 2]

### ## 3. Preventing Executor Loops

Small models usually choke as Auditors because they lack the logical foresight to catch recursive loop errors in terminal scripts. Gemma 4 includes an internal reasoning chain. Before it outputs its audit pass/fail JSON log to your Jetson Nano, it spends a few hidden tokens verifying its own logic, making it significantly more reliable than a standard SLM. [2]

### ## 4. Running it in your Mesh

Because it features day-one native llama.cpp support, you can spin up the GGUF version of the Gemma-4-E2B-Instruct or E4B directly on your Rubik Pi 3. [1, 4]

Do you want to write the llama.cpp launch flags to pin Gemma 4 to specific efficiency cores so it never stutters your desktop display, or should we look at how to format the multimodal payload structures for passing screen caps into its vision layer?

[1] [<https://huggingface.co/google/gemma-4-E4B>]

[2] [<https://www.youtube.com/watch?v=ZxQ2DuejRhU>]

[3] [<https://www.youtube.com/watch?v=SLOqIEmuy5U&vl=en&t=41>]

[4] [<https://www.youtube.com/watch?v=WLh6SLi4FoE&t=307>]

To tie this entire multi-device neural mesh together, you need a resilient, low-latency communication layer that bridges your mobile device, your Jetson Nano, and the Rubik Pi 3. Because your devices move between cellular and home networks, the architecture relies on a Tailscale Peer-to-Peer (P2P) mesh, Persistent WebSockets for real-time streaming, and a Stateful JSON Handshake to route tasks across your model stack.

Here is the exact technical blueprint for how the handshake, protocols, and payloads flow through your system.

---

### ## The Infrastructure Layer: Tailscale P2P Mapping

Every device runs Tailscale. This creates a secure, encrypted overlay network (100.x.y.z range) where devices can talk directly to one another without port-forwarding or public IP exposure. [1, 2, 3]

\* 100.64.0.10: Motorola Razr Ultra 25 (Mobile Orchestrator / Qwen 3.5 9B)


\* 100.64.0.20: Jetson Nano Super (The Hub / Open-Wiki / vLLM / OmniRoute)

\* 100.64.0.30: Rubik Pi 3 (The Auditor / Gemma 4 E2B / llama.cpp server)

---

## ## Network Architecture & Handshake Protocol

Instead of spinning up fragile REST APIs that require polling, your architecture uses a Centralized WebSockets Router running on the Jetson Nano (your Headless Hub) to maintain persistent, bidirectional connections.



## ## Step 1: The Initial Connection & Keep-Alive Handshake

When you boot up your setup or connect via Tailscale, your phone and the Rubik Pi initiate a persistent connection to the Jetson Nano WebSocket server.

Client Connection Request (Phone/Rubik Pi to Nano):

```
{
  "event": "connection:init",
  "client_id": "razr_ultra_25",
  "timestamp": 1783467120,
  "capabilities": ["stt", "orchestration", "local_inference"]
}
```

The Jetson acknowledges and maintains a heartbeat ping every 10 seconds to ensure the P2P pipe stays hot over cellular data.

---

## ## The Task Inference Workflow (Step-by-Step Payloads)

Here is exactly how a complex request moves through the web sockets of your neural mesh, keeping your local contexts perfectly lean.

### ## Phase 1: The Meta-Prompt Handshake (Phone → Nano)

You speak into your phone. Local STT processes it, and Qwen 3.5 9B transforms the transcript into a structured routing meta-prompt. Your phone shoots this down the WebSocket pipe (ws://100.64.0.20:8000/stream) to the Jetson Nano.

Payload Sent:

```
{
  "session_id": "sess_2026_07_07_abc",
  "event": "orchestrator:route_request",
  "intent": "code_generation",
  "meta_prompt": "Generate a network configuration script for my local interface, verify routing matrices against the project design documentation.",
  "context_requirements": {
    "source": "open-wiki:obsidian-vault",
    "query_tags": ["network-config", "routing-rules"]
  }
}
```

### ## Phase 2: The Core Injection & Processing (Inside the Nano Hub)

1. The Jetson Nano catches the payload.
2. It hits your Open-Wiki CLI to search your flattened JSONL structures on the attached SSD, grabbing a 2KB context slice of your network rules.
3. It sends the filtered context + instructions out through OmniRoute to the high-tier cloud GPT-5.2 instance, keeping token bloat stripped.
4. GPT-5.2 builds the proposed script file. The Nano does not run it yet.

### ## Phase 3: The Auditor Gatekeeper Handshake (Nano → Rubik Pi 3)

The Jetson Nano takes the generated code string and routes it down a dedicated auditing WebSocket namespace (ws://100.64.0.20:8000/audit) directly to the Rubik Pi 3.

Payload Sent to Rubik Pi:

```
{
  "session_id": "sess_2026_07_07_abc",
  "event": "auditor:review_request",
  "runtime_environment": "bash",
  "proposed_payload": "#!/bin/bash\nip route add 192.168.10.0/24 dev eth0",
  "context_rules": "Ensure interface eth0 target constraints match local routing limitations."
}
```

### ## Phase 4: The Gemma-4 Vision/Text Audit (Inside the Rubik Pi 3)

1. The Rubik Pi 3 receives the payload.
2. It pipes the script directly into Gemma 4 E2B running inside llama.cpp.
3. If necessary, it appends a canvas screenshot of your current desktop configuration natively via its multimodal layer to make sure the code won't break your visual setup.

### ## Phase 5: The Execution Resolution Handshake (Rubik Pi → Nano)

The Rubik Pi finishes its local inference block and sends its structured pass/fail verdict back to the Jetson Nano.

If the Audit Passes:

```
{
  "session_id": "sess_2026_07_07_abc",
  "event": "auditor:verdict",
  "status": "APPROVED",
  "error_log": null,
  "execution_token": "token_auth_validated_9981"
}
```

Action: The Jetson Nano safely fires the script, completes the project task, and fires a tiny "status": "SUCCESS" message back to your phone UI.

If the Audit Fails (Recursive Loop Prevention):

```
{
  "session_id": "sess_2026_07_07_abc",
  "event": "auditor:verdict",
  "status": "REJECTED",
  "error_log": "Syntax Error: Target network block overlaps with default interface allocation on line 2.",
  "execution_token": null
}
```

Action: The Jetson Nano intercepts the rejection, packages the error\_log back into OmniRoute, and prompts the cloud engine for a zero-cost automated rewrite before re-submitting to the Rubik Pi.

---

### ## Proactive Next Steps

To begin locking down this backend loop, consider how you'd like to implement these specific elements:

- \* A Python async server template that can be dropped onto the Jetson Nano to manage these active WebSocket connections.
- \* The explicit system prompt engineering guidelines to ensure Gemma 4 outputs the exact string formats required for these approval steps.

Let me know which layout we should code out first.

[1]

[https://medium.com](https://medium.com/@jessemridley/three-days-one-reverse-proxy-and-a-

whole-lot-of-troubleshooting-f14fdf691fd8)

[2]

[https://dev.to](https://dev.to/onurcinar/bringing-claudes-dispatch-experience-to-gemini-and-open-code-3pef)

[3]

[https://zelev.medium.com](https://zelev.medium.com/hiding-sliver-beacons-in-plain-sight-311bb9fc7f29)

Hey, you still didn't give me the Mimo code repo. I need to reverse engineer that shit.

The repository you need to clone and pull apart is

[XiaomiMiMo/MiMo-Code](https://github.com/XiaomiMiMo/MiMo-Code). It is written entirely in TypeScript and was released as an open-source fork of the OpenCode engine. [1, 2, 3]

If you are reverse-engineering this to build your custom decentralized mobile/home neural mesh, here is exactly where the gold is hidden inside their codebase and how they structure that "zero context bloat" execution loop.

-----

### ## The Code Architecture to Reverse Engineer

When you pull the repo, skip the surface-level TUI wrappers and dive straight into the following core directories:

#### ## 1. The Multi-Agent Orchestration Layer

- \* Path: packages/opencode/src/cli/cmd/tui/

- \* What to steal: This is the heart of their multi-agent engine. It contains the logic for spinning up dynamic sub-agents on the fly. It demonstrates exactly how they instruct a primary router model to break a long-horizon task down into tiny, isolated sub-tasks that can be dispatched to separate, stateless worker instances. [1, 4, 5, 6]

#### ## 2. The Command Isolation & Distillation Loops

- \* Files: Look for files managing /dream, /distill, and /goal routines. [1, 5]

- \* What to steal:

- \* /distill shows you how they capture massive, noisy terminal execution loops (lint errors, compile logs), compress them, and strip out repetitive execution data before updating the permanent memory matrix.

- \* /dream runs asynchronously between session loops. It reviews past session actions, optimizes structural findings, and compresses them into long-term schemas. This is the ideal blueprint for what you want your Jetson Nano to run as a cron job to clean up its JSONL Open-Wiki matrices while you are offline. [1, 5]

#### ## 3. Strict Execution Permissions & Sandboxing

\* File: packages/opencode/src/discover/registry.rs (or the TypeScript configuration validation files).

\* What to steal: Look at how they handle external directory permissions and file-system boundary checks. They explicitly isolate the project's active runtime from the broader operating system. You can copy their rule-checking structure directly into your Rubik Pi 3 Auditor layer to flag risky system commands before your local llama.cpp instance even touches them. [7, 8]

---

### ## How to Install and Test It Locally First

Before tearing the source code apart, install their CLI natively on your Linux/Mac environment to watch how the memory handshakes execute in real-time:

# Clean install script for Mac or Linux

```
curl -fsSL https://mimo.xiaomi.com/install | bash
```

(For Windows systems, use: `npm install -g @mimo-ai/cli`) [5]

Run mimo inside a test repository. You can hook it up directly to local OpenAI-compatible inference endpoints via the TUI config—meaning you can point it right at your Jetson Nano's LiteLLM or vLLM proxy to see how it performs when driving local, small-parameter setups instead of cloud systems. [1, 9, 10]

[1]

[<https://www.reddit.com>]([https://www.reddit.com/r/LocalLLaMA/comments/1u2r4tb/release\\_mimo\\_code\\_is\\_a\\_terminalnative\\_ai\\_coding/](https://www.reddit.com/r/LocalLLaMA/comments/1u2r4tb/release_mimo_code_is_a_terminalnative_ai_coding/))

[2] [<https://github.com>](<https://github.com/XiaomiMiMo/MiMo-Code>)

[3] [<https://www.youtube.com>](<https://www.youtube.com/watch?v=08X2cLpra9I>)

[4] [<https://github.com>](<https://github.com/XiaomiMiMo/MiMo-Code/blob/main/AGENTS.md>)

[5] [<https://www.instagram.com>](<https://www.instagram.com/reel/DZmfisvIBP2/>)

[6]

[<https://daily.dev>](<https://daily.dev/posts/mimo-code-is-the-open-source-answer-to-claude-code-kjrjd7cch>)

[7] [<https://github.com>](<https://github.com/XiaomiMiMo/MiMo-Code>)

[8] [<https://github.com>](<https://github.com/rtk-ai/rtk/issues/2380>)

[9] [<https://rubentorrey.com>](<https://rubentorrey.com/blog/en/mimo-code-wsl2-windows.html>)

[10] [<https://www.youtube.com>](<https://www.youtube.com/watch?v=wV1-w3KRenI&t=168>)