



namhokaist/

appgen-qwen3-model-e-simplepath-sparse-s10warm-v42-h200x8-s100-20260731memfix-step60

like 0



Image-Text-to-Text



Transformers



Safetensors



luca0621/appgen-sft-ngc-v1



Yuxiang007/AMEX

qwen3_vl

qwen3-vl

android

gui-agent

sft

appgen

conversational



License: apache-2.0



Deploy

→ Copy to bucket

NEW

Use this model



Model card



Files



xet



Community

Downloads last month

-



Safetensors

Model size

9B params

Tensor type

BF16



Chat template



Files info



Inference Providers

NEW



Image-Text-to-Text

This model isn't deployed by any Inference Provider.



Ask for provider support



Model tree for namhokaist/appgen-qwen3-model-e-simplepath-sparse-s10warm-...

Base model

[Qwen/Qwen3-VL-8B-Instruct](#)



Finetuned (447) [this model](#)



Datasets used to train namhokaist/appgen-qwen3-model-e-simplepath-sparse-s1...



Yuxiang007/AMEX

Updated Sep 27, 2024 •  1.28k •  30

 luca0621/appgen-sft-ngc-v1

 Viewer • Updated 11 days ago •  7.18k •  67

 Edit model card

AppGen Qwen3-VL frozen SFT — A-variant arm E

This is arm ngc_anchor_lr2p5e7 from a preregistered four-arm follow-up to the best NGC configuration A. The language model was fully fine-tuned while the complete Qwen3-VL visual tower, merger, and deep-stack mergers remained frozen. Publication verifies every model.visual tensor byte-for-byte against the pinned base and verifies that language-model weights changed.

Training contract

- Base: Qwen/Qwen3-VL-8B-Instruct at commit 0c351dd01ed87e9c1b53cbc748cba10e6187ff3b
- Sources: luca0621/appgen-sft-ngc-v1 at 769ea99dbc4ff190048ae0db37eb6310dba595e0 and Yuxiang007/AMEX at 17196b29c88dd48a7fb90ef9131bc5c7bf39f26e
- Dataset variant: ngc_anchor
- Dataset SHA-256: 8890e2fe596e4d78714ce4309f2a59ed9632dd5d656f758c8dc4608d46a9c6bc
- Exposures: 3,588; unique semantic examples: 3,168
- Direct-grounding exposures: 400
- Completion-retention replay exposures: 420
- Image provenance: synthetic AppGen HTML-to-PNG
- Coordinates: normalized 0–1000

- Prompt: proven A prompt; SHA-256
67ff8adb0e78a617f3d0edcf196d4e4cc3239967a8c30619fc5afc14484ee8c0
- Optimizer: full-language AdamW, learning rate 2.5e-07, cosine schedule, 5% warmup
- Batch: microbatch 4 × two GPUs × gradient accumulation 4 = global batch 32
- Epochs: 1; optimizer updates: 113
- Loss scale: Swift default
- Frozen: visual encoder and aligner; trainable: language model and LM head

Only the final checkpoint (checkpoint-113) is published. The exact system prompt is `appgen_system_prompt.txt`; `run_manifest.json` records exact hashes, the dataset receipt, overlap audit, and weight verification. Intermediate 25-step checkpoints remain local.

Evaluation and limitations

The training/evaluation overlap gate compares EXIF-transposed decoded RGBA pixels plus image dimensions against the seven pinned AW7 test suites. It also compares normalized instruction hashes. The exact pinned receipt must report zero train/eval pixel and instruction overlaps for every data variant and suite before publication. Static grounding benchmarks are not a substitute for interactive AndroidWorld task success, and no score is claimed in this card.

This model is for Android visual-agent research, not safety-critical autonomous deployment.

TOS

Privacy

About

Careers

Website

Models

Datasets

Spaces

Pricing

Docs

