


[Home](#) > [All Models](#) > **Qwen3.5-2B**

Qwen3.5-2B

Balanced 2B parameter model from the Qwen3.5 series for on-device inference.

Qwen3.5 is the latest multilingual language model series from Alibaba Cloud with improved reasoning and instruction-following capabilities over Qwen3.

[Quick Start](#)[View README](#)

q4_0

 GenieX - llama.cpp



Qwen3.5-2B



Snapdragon X2 Elite CRD Snapdragon® X2 Elite



NPU ▾

CONTEXT LENGTH

4096 ▾

tokens

PREFILLING SPEED

894

tokens/s

DECODING SPEED

24.6

tokens/s

* This model uses the llama.cpp community runtime. [Request optimization](#) if the performance doesn't meet expectation.

Quick Start

Windows

Android

Linux

① Install Windows CLI App

[Download for Windows ARM](#)

② Run the Model ⓘ

Paste into CLI and run the following code.

 GenieX - llama.cpp

```
geniex infer unsloth/Qwen3.5-2B-GGUF
```



Or [download the model](#).

ⓘ For application and server integration, [see Docs](#).



Model Repository



Hugging Face



Research Paper



Technical Details

model_infer_id: unsloth/Qwen3.5-2B-GGUF

Quantization Type: Q4_0 (GGUF)

Model Size: 1.21 GB

Applicable Scenarios

✓ Dialogue

✓ Content Generation

Supported Form Factors



Phone



Tablet



IoT

License

Model: [APACHE-2.0](#)

Terms of Use: [Qualcomm® Generative AI usage and limitations](#)

Tags

llm ⓘ

vlm ⓘ

generative-ai ⓘ

Supported Devices

 Dragonwing IQ-9075 EVK

 Snapdragon 8 Elite Gen 5 QRD

 Snapdragon 8 Elite QRD

 Snapdragon X Elite CRD

 Snapdragon X2 Elite CRD

Supported Chipsets

 Qualcomm® Dragonwing™ IQ-9075

 Snapdragon® 8 Elite Mobile

 Snapdragon® 8 Elite Gen 5 Mobile

 Snapdragon® X Elite

 Snapdragon® X2 Elite

Related Models

[See all models >](#)



Floyd

What's the recommended dosage for this medicine?




AI Assistant

Analyzing...



Qwen3.5-0.8B



Qwen3-1.7B

Ultra-compact 0.8B parameter model
from the Qwen2.5 series for edge

Multilingual 1.7B parameter language
model excelling in reasoning and

Looking for more? See models created by industry leaders.

[Discover Model Makers](#)



[Terms of Use](#) [Privacy Policy](#) [Cookie Policy](#) [Contact Us](#)

© 2026 Qualcomm Technologies, Inc. and/or its affiliated companies.

References to "Qualcomm" may mean Qualcomm Incorporated, or subsidiaries or business units within the Qualcomm corporate structure, as applicable.

Qualcomm Incorporated includes our licensing business, QTL, and the vast majority of our patent portfolio. Qualcomm Technologies, Inc., a subsidiary of Qualcomm Incorporated, operates, along with its subsidiaries, substantially all of our engineering, research and development functions, and substantially all of our products and services businesses, including our QCT semiconductor business.

Materials that are as of a specific date, including but not limited to press releases, presentations, blog posts and webcasts, may have been superseded by subsequent events or disclosures.

Nothing in these materials is an offer to sell any of the components or devices referenced herein.