



qualcomm/Qwen3.5-0.8B



like 0

Follow Qualcomm 1.09k



Text Generation



PyTorch

llm

vlm

generative_ai

android



License: other



→ Copy to bucket NEW



Model card



Files



xet



Community

Downloads last month

-

Downloads are not tracked for this model. [How to track](#)



⚡ Inference Providers NEW



Text Generation

This model isn't deployed by any Inference Provider.



Ask for provider support



Edit model card



Floyd

What's the recommended dosage for this medicine?



AI Assistant

Analyzing...



[🔗](#) Qwen3.5-0.8B: Optimized for Qualcomm Devices

Qwen3.5 is the latest multilingual language model series from Alibaba Cloud with improved reasoning and instruction-following capabilities over Qwen3.

This is based on the implementation of Qwen3.5-0.8B found [here](#). This repository contains pre-exported model files optimized for Qualcomm® devices. You can use the [Qualcomm® AI Hub Models](#) library to export with custom configurations. More details on model performance across various devices, can be found [here](#).

Qualcomm AI Hub Models uses [Qualcomm AI Hub Workbench](#) to compile, profile, and evaluate this model. [Sign up](#) to run these models on a hosted Qualcomm® device.

[🔗](#) Deploying Qwen3.5-0.8B on-device

Follow the [GenieX quickstart](#) to install GenieX and deploy the model on a target device.

[🔗](#) Getting Started

There are two ways to deploy this model on your device:

Option 1: Download Pre-Exported Models

Below are pre-exported model assets ready for deployment.

Runtime	Precision	Chipset	SDK Versions	Download
GENIEX_LLAMACPP	q4_0	Universal		Download

For more device-specific assets and performance metrics, visit [Qwen3.5-0.8B on Qualcomm® AI Hub](#).

Option 2: Export with Custom Configurations

Use the [Qualcomm® AI Hub Models](#) Python library to compile and export the model with your own:

- Custom weights (e.g., fine-tuned checkpoints)
- Custom input shapes
- Target device and runtime configurations

This option is ideal if you need to customize the model beyond the default configuration provided here.

See our repository for [Qwen3.5-0.8B on GitHub](#) for usage instructions.

Model Details

Model Type: Model_use_case.text_generation

Model Stats:

- Model architecture: Transformer with GQA and SwiGLU
- Supported languages: 100+ languages and dialects

- TTFT: Time To First Token is the time it takes to generate the first response token.
This is expressed as a range because it varies based on the length of the prompt.
- Response Rate: Rate of response generation after the first response token.

 Performance Summary

Model	Runtime	Precision	Chipset	Context Length	Response Rate (tokens per second)	Time To F (range, s)
Qwen3.5-0.8B	GENIEX_LLAMACPP	q4_0	Snapdragon® 8 Elite Gen 5 Mobile	512	62.131097	0.23427574 - 0.93710299
Qwen3.5-0.8B	GENIEX_LLAMACPP	q4_0	Snapdragon® 8 Elite Gen 5 Mobile	512	50.763998	0.27429675
Qwen3.5-0.8B	GENIEX_LLAMACPP	q4_0	Snapdragon® 8 Elite Gen 5 Mobile	512	35.869292	0.11880325
Qwen3.5-0.8B	GENIEX_LLAMACPP	q4_0	Snapdragon® 8 Elite Gen 5 Mobile	4096	59.414176	0.36005078 - 11.521625
Qwen3.5-0.8B	GENIEX_LLAMACPP	q4_0	Snapdragon® 8 Elite Gen 5 Mobile	4096	56.679703	0.39175446 - 12.536143
Qwen3.5-0.8B	GENIEX_LLAMACPP	q4_0	Snapdragon® 8 Elite Gen 5 Mobile	4096	41.909392	0.11803384 - 3.777083
Qwen3.5-0.8B	GENIEX_LLAMACPP	q4_0	Snapdragon® 8 Elite Mobile	512	72.950102	0.26924750 -

Model	Runtime	Precision	Chipset	Context Length	Response Rate (tokens per second)	Time To First Token (range, seconds)
						1.07699000
Qwen3.5-0.8B	GENIEX_LLAMACPP	q4_0	Snapdragon® 8 Elite Mobile	512	62.441461	0.26454675
Qwen3.5-0.8B	GENIEX_LLAMACPP	q4_0	Snapdragon® 8 Elite Mobile	512	37.52486	0.1532055
Qwen3.5-0.8B	GENIEX_LLAMACPP	q4_0	Snapdragon® 8 Elite Mobile	4096	49.751244	0.42967528 13.749609
Qwen3.5-0.8B	GENIEX_LLAMACPP	q4_0	Snapdragon® 8 Elite Mobile	4096	48.911714	0.43504246 13.921359
Qwen3.5-0.8B	GENIEX_LLAMACPP	q4_0	Snapdragon® 8 Elite Mobile	4096	35.893754	0.15914856 5.092754
Qwen3.5-0.8B	GENIEX_LLAMACPP	q4_0	Snapdragon® X2 Elite	512	88.774407	0.09157649 - 0.36630599
Qwen3.5-0.8B	GENIEX_LLAMACPP	q4_0	Snapdragon® X2 Elite	512	87.673154	0.0942735
Qwen3.5-0.8B	GENIEX_LLAMACPP	q4_0	Snapdragon® X2 Elite	512	40.640275	0.1113045
Qwen3.5-0.8B	GENIEX_LLAMACPP	q4_0	Snapdragon® X2 Elite	4096	81.095293	0.12939821 - 4.14074300
Qwen3.5-0.8B	GENIEX_LLAMACPP	q4_0	Snapdragon® X2 Elite	4096	80.304757	0.14290928 4.573097
Qwen3.5-0.8B	GENIEX_LLAMACPP	q4_0	Snapdragon® X2 Elite	4096	39.613582	0.11075478 -

Model	Runtime	Precision	Chipset	Context Length	Response Rate (tokens per second)	Time To First Token (range, seconds)
Qwen3.5-0.8B	GENIEX_LLAMACPP	q4_0	Qualcomm® Dragonwing™ IQ-X7181	512	63.044187	0.12997225
Qwen3.5-0.8B	GENIEX_LLAMACPP	q4_0	Qualcomm® Dragonwing™ IQ-X7181	512	29.988464	0.2527895
Qwen3.5-0.8B	GENIEX_LLAMACPP	q4_0	Qualcomm® Dragonwing™ Q-8750	512	72.950102	0.26924750 - 1.07699000
Qwen3.5-0.8B	GENIEX_LLAMACPP	q4_0	Qualcomm® Dragonwing™ Q-8750	512	62.441461	0.26454675
Qwen3.5-0.8B	GENIEX_LLAMACPP	q4_0	Qualcomm® Dragonwing™ Q-8750	512	37.52486	0.1532055
Qwen3.5-0.8B	GENIEX_LLAMACPP	q4_0	Qualcomm® Dragonwing™ IQ-9075	4096	38.271652	0.46651725 14.928552
Qwen3.5-0.8B	GENIEX_LLAMACPP	q4_0	Qualcomm® Dragonwing™ IQ-9075	4096	39.346842	0.46202675 14.784856
Qwen3.5-0.8B	GENIEX_LLAMACPP	q4_0	Qualcomm® Dragonwing™ IQ-9075	4096	14.03834	0.33502103 - 10.7206730
Qwen3.5-0.8B	GENIEX_LLAMACPP	q4_0	Qualcomm® Dragonwing™ IQ-X7181	4096	51.537779	0.21393696 6.845983

Model	Runtime	Precision	Chipset	Context Length	Response Rate (tokens per second)	Time To First Token (range, seconds)
Qwen3.5-0.8B	GENIEX_LLAMACPP	q4_0	Qualcomm® Dragonwing™ IQ-X7181	4096	48.821458	0.23768365 7.605877
Qwen3.5-0.8B	GENIEX_LLAMACPP	q4_0	Qualcomm® Dragonwing™ IQ-X7181	4096	28.38794	0.25720621 8.230599
Qwen3.5-0.8B	GENIEX_LLAMACPP	q4_0	Qualcomm® Dragonwing™ Q-8750	4096	49.751244	0.42967528 13.749609
Qwen3.5-0.8B	GENIEX_LLAMACPP	q4_0	Qualcomm® Dragonwing™ Q-8750	4096	48.911714	0.43504246 13.921359
Qwen3.5-0.8B	GENIEX_LLAMACPP	q4_0	Qualcomm® Dragonwing™ Q-8750	4096	35.893754	0.15914856 5.092754

License

- The license for the original implementation of Qwen3.5-0.8B can be found [here](#).

References

- [Qwen3.5: Towards Native Multimodal Agents](#)
- [Source Model Implementation](#)

Community

- Join [our AI Hub Slack community](#) to collaborate, post questions and learn more about on-device AI.
- For questions or feedback please [reach out to us](#).

Usage and Limitations

This model may not be used for or in connection with any of the following applications:

- Accessing essential private and public services and benefits;
- Administration of justice and democratic processes;
- Assessing or recognizing the emotional state of a person;
- Biometric and biometrics-based systems, including categorization of persons based on sensitive characteristics;
- Education and vocational training;
- Employment and workers management;
- Exploitation of the vulnerabilities of persons resulting in harmful behavior;
- General purpose social scoring;
- Law enforcement;
- Management and operation of critical infrastructure;
- Migration, asylum and border control management;
- Predictive policing;
- Real-time remote biometric identification in public spaces;
- Recommender systems of social media platforms;
- Scraping of facial images (from the internet or otherwise); and/or
- Subliminal manipulation

 System theme

Company

TOS

Privacy

About

Careers

Website

Models

Datasets

Spaces

Pricing

Docs

