

Failure Log

Review weekly. Patterns emerge fast. Every entry becomes a test case (Skill 2) and feeds the challenge dataset (PDF Step 2).

Entry Format

[DATE] — [SYSTEM/PROJECT]

Symptom: *What happened? What did you observe?*

Expected Behavior: *What should have happened?*

Root Cause: *What actually went wrong at the source?*

Failure Type (pick one):

- ☐ Context Degradation — quality dropped as session/context got long
- ☐ Specification Drift — agent/system forgot or drifted from the original intent
- ☐ Sycophantic Confirmation — agent confirmed incorrect data and built on it
- ☐ Tool Selection Error — wrong tool picked, wrong approach taken by agent
- ☐ Cascading Failure — one failure propagated through the chain, no correction
- ☐ Silent Failure — output looked correct but was functionally wrong

Blast Radius: *Who/what was affected? How bad?*

- Severity: Low / Medium / High / Critical
- Reversible: Yes / No / Partial
- Frequency: One-off / Recurring / Systemic

Fix Applied: *What did you do to resolve it?*

Prevention: *What would stop this from happening again? Config change? Guardrail? Test case?*

Test Case Generated: *Write one. Input → Expected Output → Pass/Fail criteria.*

Input:

Expected Output:

Pass/Fail:

Weekly Pattern Review

Week	Most Common Type	Recurring System	Action Taken

Failure Type Frequency Tracker

Type	Count	Last Seen	Trend
Context Degradation			
Specification Drift			
Sycophantic Confirmation			
Tool Selection Error			
Cascading Failure			
Silent Failure			