

Here is the complete, unified master specification document. It consolidates every section from the initial pilot draft, the multi-model cross-auditing and RLVR pipelines, the universal skill/tool extraction engine, the distributed RAG manifest schemas, the OpenWiki TUI / wiki_md/ protocol, and the full scaffolding script into a single, fully realized document.

SYSTEM SPECIFICATION: COGNITIVE REPOSITORY ARCHITECTURE

Document Target: Autonomous Planning, Auditing, and Execution Agents Role: Injected Structural, Operational & Synchronization Protocol Repository Model: 5+1 Tier Cognitive Memory, OpenWiki TUI File Management, Multi-Model Cross-Auditing, Sandboxed Red Gatekeeping, and Distributed RAG Manifests Version: 3.0 (Definitive Master Specification)

1. Core Architectural Invariants

As the planning agent, you must enforce the separation of cognitive memory layers across the filesystem. Adhere strictly to the following invariants:

1. Sensory Immutability (Tier 1): Never modify, overwrite, or delete assets in 01_raw_sources/. All raw inputs are read-only sources of truth.
2. Human-Machine Decoupling (Tier 2 vs. Tier 3): Human-readable semantic knowledge lives in 02_wiki_md/ (.md). Machine-readable retrieval chunks live in 03_recall_cache/ (.jsonl, vectors). Never dump raw tabular JSONL or vector embeddings into the wiki layer.
3. Procedural vs. Episodic Separation (Tier 4 vs. Tier 5): Executable skills, scripts, and policies belong in 04_skills_runtime/. Trajectory traces, error logs, and verifier evaluations belong in 05_episodic_logs/. Never place runtime logic inside logging directories or vice versa.
4. Metacognitive Priming: Before initiating multi-step tasks, read MAP.md or parse the root manifest.jsonl to resolve entity references rather than performing blind recursive scans across the filesystem.
5. LLM Wiki Protocol (wiki_md/): The semantic vault operates as an organization-wide wiki standard governed directly by the Files Executive Agent via the OpenWiki TUI.
6. Dual Extraction Mandate: Every document ingested into the ecosystem must be systematically screened for both Procedural Skills (prompt templates, behavioral heuristics) and Executable Tools (scripts, wrappers, deterministic CLI invocations).
7. Red Auditor Gatekeeping: No episodic execution trace or generated script may enter the recursive training stream without passing an isolated sandbox evaluation by the Red Auditor Agent.
8. Universal Source of Truth: All local auditor nodes and the sandboxed Red Auditor must evaluate claims and script outcomes against the shared master memory bank and database.
9. Distributed Manifest Coverage: Every repository root and major structural subsystem must maintain its own local manifest.jsonl to power low-overhead hierarchical RAG retrieval.

2. Directory Schema & Access Control Matrix

 vault_root/	
—  MAP.md	# Master human-readable index & ontology graph
(Read-Heavy / Agent Update)	
—  manifest.jsonl	# Root-level RAG registry & global hash table (Append / Sync)
—  01_raw_sources/	# TIER 1: COLD SENSORY ARCHIVE (Access:

READ-ONLY)

- manifest.jsonl # Local RAG catalog of all immutable raw sources
- pdf/ # Source whitepapers, specs, data sheets (*.pdf)
- media/ # Video captures, audio memos, transcripts (*.mp4, *.wav, *.png)
- text/ # Raw text scrapes, dumps, API payload exports (*.txt, *.html)

02_wiki_md/ # TIER 2: THE LLM WIKI LAYER (Access: READ /

WRITE - OpenWiki TUI)

- manifest.jsonl # Local RAG catalog of conceptual nodes & graph links
- concepts/ # Atomic linked notes, theory, and domain models (*.md)
- architectures/ # System blueprints, data flows, interface specs (*.md)
- entities/ # Registries, schema contracts, hardware profiles (*.md)
- indexes/ # Maps of Content (MOCs) and taxonomy clusters (*.md)

03_recall_cache/ # TIER 3: WORKING MEMORY ACCELERATOR

(Access: REBUILD / OVERWRITE)

- manifest.jsonl # Local RAG catalog of vector & chunk shards
- jsonl/ # Pre-tokenized passages for high-speed prompt injection
- vectors/ # Dense vector indices (*.bin, *.faiss, *.hnsf checkpoints)
- kv_store/ # Low-latency key-value entity lookups (*.db, *.json)

04_skills_runtime/ # TIER 4: PROCEDURAL REPERTOIRE & TOOLS

(Access: VERSION-CONTROLLED)

- manifest.jsonl # Local RAG catalog of all extracted skills & tools
- prompt_skills/ # Modular SKILL.md prompt definitions (*.md, *.yaml)
- extracted_tools/ # Deterministic tools mined from ingested

documentation

- cli/ # Extracted command-line utilities (*.sh)
- wrappers/ # Extracted API interfaces & micro-functions (*.py)
- runtimes/ # Master operational scripts and execution hooks (*.py, *.sh)
- policies/ # Validation guards, retry policies, schemas (*.json, *.yaml)

05_episodic_logs/ # TIER 5: TELEMETRY & EPISODIC RUNS (Access:

APPEND-ONLY)

- manifest.jsonl # Local RAG catalog of sessions, audits, and verifiers
- daily_driver_sync/ # Ingested logs from edge device sessions (*.jsonl)
- trajectories/ # Multi-model traces: Query, Executor, Frontier (*.jsonl)
- red_audit_sandbox/ # Red Auditor evaluation reports & pass/fail quarantine

(*.jsonl)

—  rlvr_verifiers/	# Graded rewards, assertion outcomes (+1.0 / -1.0) (*.jsonl)
—  hygiene_reports/	# Schema integrity audits, broken link checks (*.md)

3. Tier Operational Specifications

Root: Metacognitive Index (MAP.md & manifest.jsonl)

* MAP.md: Must maintain top-level navigation links using Markdown cross-references to key notes in 02_wiki_md/ and entry points in 04_skills_runtime/.

* manifest.jsonl (Root): Master catalog mapping file paths, tiers, cryptographic checksums, and cross-tier linkages:

```
{"id": "SRC-0042", "path": "01_raw_sources/pdf/spec_v1.pdf", "tier": 1, "sha256": "e3b0c44298fc1c149afb4c8996fb92427ae41e4649b934ca495991b7852b855", "last_modified": "2026-09-03T17:00:00Z", "linked_wiki": "02_wiki_md/architectures/spec_v1.md"}
```

Tier 1: 01_raw_sources/ (Cold Archive & Sensory Buffer)

- * Write Policy: Write-once upon raw intake. Never mutate or reformat existing files.
- * Naming Convention: YYYYMMDD_source_title.[ext] (lowercase, snake_case).
- * Scope: Source PDFs, audio transcripts, sensor captures, video recordings, raw scrapes.

Tier 2: 02_wiki_md/ (The LLM Wiki Layer)

- * Write Policy: Managed via OpenWiki TUI by the Files Executive Agent.
- * Format: Pure Markdown with strict YAML frontmatter:

```
---
id: wiki_rag_indexing_protocol
tier: 2
type: architecture
created: 2026-09-03
updated: 2026-09-03
author: files_executive_agent
sources:
  - "01_raw_sources/pdf/rag_system_spec.pdf"
extracted_skills:
  - "04_skills_runtime/prompt_skills/manifest_sync.md"
extracted_tools:
  - "04_skills_runtime/extracted_tools/cli/build_manifest.py"
tags:
  - llm_wiki
  - openwiki
  - indexing
---
```

RAG Indexing Protocol

Context & Definition

Structural mechanism for keeping the LLM Wiki synchronized with high-speed caches...

Architectural Interfaces

The `[[manifest_registry_spec]]` defines how this node is indexed by the `[[files_executive_agent]]`.

References

- [[hierarchical_manifest_routing]]
- [[procedural_tool_extraction]]

Tier 3: 03_recall_cache/ (Working Memory Accelerator)

* Write Policy: Agent-generated during build or indexing steps. May be deleted and rebuilt safely.

* Chunk Schema (03_recall_cache/jsonl/*.jsonl):

```
{
  "chunk_id": "CHK-8901",
  "parent_doc": "02_wiki_md/concepts/rlvr.md",
  "tokens": 256,
  "content": "RLVR verifiers score execution traces against deterministic unit tests...",
  "metadata": {
    "tier": 2,
    "topic": "rlvr"
  }
}
```

Tier 4: 04 skills runtime/ (Procedural Memory: Skills & Tools)

* Write Policy: Version-controlled tooling, prompts, and instructions.

* Prompt Skill Format (prompt_skills/): Standard modular prompt format containing purpose, required inputs, tool dependencies, and output schemas (SKILL.md).

* Extracted Tools (extracted_tools/): Standalone executable code blocks extracted from documentation. Non-interactive scripts with standard exit codes (0 = success, non-zero = error).

Tier 5: 05 episodic logs/ (Episodic Memory & Telemetry)

* Write Policy: Append-only during agent task execution.

* Trajectory Schema (trajectories/YYYYMMDD_session.jsonl):


```
{
  "timestamp": "2026-09-03T17:01:35Z",
  "step": 1,
  "task_id": "TASK-104",
  "prompt_hash": "a1b2c3",
  "tool_call": "run_linter",
  "exit_code": 0,
  "response_snippet": "OK"
}
```


* RLVR Log Schema (rlvr_verifiers/YYYYMMDD_eval.jsonl):

```
{
  "timestamp": "2026-09-03T17:01:40Z",
  "task_id": "TASK-104",
  "verifier_id": "syntax_test",
  "reward": 1.0,
  "feedback": "All assertions passed."
}
```


4. The LLM Wiki Protocol & OpenWiki Files Executive Agent

The LLM Wiki (02_wiki_md/) serves as the persistent organizational spine. The Files Executive Agent manages it through the OpenWiki TUI:





1. TUI-Driven Lifecycle Management: The Files Executive Agent runs inside the OpenWiki TUI, monitoring file additions, moves, renames, and refactors across all tiers.
 2. Deterministic Linking & Backlinks: Automatically verifies bidirectional Markdown links ([[entity_name]]), ensuring the LLM Wiki remains an intact graph without orphaned concepts.
 3. Extraction Intake Coordination: Acts as the bridge between incoming raw sources in 01_raw_sources/, synthesizing structured knowledge into 02_wiki_md/, and passing extracted procedural artifacts to 04_skills_runtime/.
 4. Manifest Synchronization: Whenever the agent modifies a node in wiki_md/, it immediately writes updated hashes, token counts, and entity tags to both 02_wiki_md/manifest.jsonl and the root manifest.jsonl.
 5. Universal Skill & Tool Extraction Engine
- The compilation agent processes incoming documentation through a dual-channel extraction filter:



[prompt_skills/*.md] [extracted_tools/*]

Extraction Classification Rules

1. Targeting Skills: If a passage describes how to reason, how to structure an analysis, or how to coordinate a sequence of tasks, format it as a markdown skill with YAML frontmatter into 04_skills_runtime/prompt_skills/.

2. Targeting Tools: If a passage contains executable code, CLI command flags, API request formats, or data transformation utilities, extract it into an executable script (.py or .sh) with explicit argument parsing into 04_skills_runtime/extracted_tools/.

3. Registration: Update the local 04_skills_runtime/manifest.jsonl immediately upon extraction.

6. Hierarchical manifest.jsonl Specification for Distributed RAG

Every directory root must hold a dedicated manifest.jsonl file. This prevents full-vault scanning during localized RAG operations.

JSONL Schema Structure

Each line in any manifest.jsonl must strictly adhere to this schema:

```
{
  "id": "UUID-OR-PATH-HASH",
  "path": "relative/path/to/file.ext",
  "tier": 4,
  "category": "extracted_tool",
  "sha256": "e3b0c44298fc1c149afb4c8996fb92427ae41e4649b934ca495991b7852b855",
  "tokens": 412,
  "semantic_summary": "CLI utility to extract audio streams from MP4 video containers.",
  "entities_extracted": ["ffmpeg", "audio_processing", "mp4_to_wav"],
  "skills_tools_extracted": [
    {"type": "tool", "name": "extract_audio_stream", "path":
"04_skills_runtime/extracted_tools/cli/extract_audio.sh"},
    {"type": "skill", "name": "audio_preprocessing_policy", "path":
"04_skills_runtime/prompt_skills/audio_prep.md"}
  ],
  "provenance_source": "01_raw_sources/pdf/media_processing_guide.pdf",
  "last_synced": "2026-09-03T17:05:00Z"
}
```

RAG Query Traversal Protocol


1. Directory-Level Scope Matching: The RAG agent inspects root manifest.jsonl to identify candidate subdirectories.

2. Localized Search: The agent reads only the relevant sub-tier manifest.jsonl (e.g., 04_skills_runtime/manifest.jsonl), matching semantic summaries and entities.

3. Targeted Ingestion: Only the exact chunk or tool indicated by the manifest record is loaded into active Working Memory.

7. End-of-Day P2P Sync, Cross-Auditor & Sandboxed Red Auditor Workflow

When the daily driver reconnects to the home environment via the peer-to-peer server mesh, the Home Script Collector / Auditor Agent executes the cross-audit:



Tri-Model Trace Alignment

1. Query Model: Captures initial intent decomposition, system prompts, and task planning.

2. Executor Model: Captures deterministic tool invocations, bash runs, and code snippets.

3. Frontier Model: Captures fallback escalations, complex conceptual synthesis, and recovery reasoning.

Sandbox Evaluation & Dispositions

* Determinism: The extracted script must execute in isolation without hidden dependencies.

* Grounding Authority: Must agree with the universal memory bank without invented parameters.

* Dispositions:

* VERDICT_APPROVED (+1.0): Moved to upload queue for recursive training.

* VERDICT_REJECTED (-1.0): Logged to 05_episodic_logs/red_audit_sandbox/ with an explicit failure reason, triggering recursive self-correction in Tier 4 policies.


8. Cloud Scriptor & RLVR Recursive Training Pipeline

[Approved Trajectories]




Cloud Pipeline Intake: GCS Bucket
gs://<repo>-rlvr-training-pipeline/

Triggers Cloud Function / Cloud Run




Downstream Artifact Sync



| - Verified Run Logs —▶ 05_episodic_logs/rlvr_verifiers/ |

9. Agent Execution & Planning Protocols



Operational Steps for the Planning Agent

Step 1: Locating Information (Read Pass)

1. When asked a domain question, check 03_recall_cache/ first for relevant chunks.
2. If full context is required, resolve the source path via MAP.md and read the relevant note in 02_wiki_md/.
3. Only query 01_raw_sources/ if source verification or citation check is explicitly requested.

Step 2: Scaffolding or Modifying Repo Content (Write Pass)

1. When processing new source material:
 - * Subsection 1: Place the original asset into 01_raw_sources/<type>/.

- * Subsection 2: Screen document and synthesize conceptual notes into 02_wiki_md/<category>/ via OpenWiki TUI conventions.
- * Subsection 3: Extract procedural skills to 04_skills_runtime/prompt_skills/ and executable tools to 04_skills_runtime/extracted_tools/.
- * Subsection 4: Regenerate the corresponding search chunks in 03_recall_cache/jsonl/, update the sub-tier manifest.jsonl, and append the record to the root manifest.jsonl and MAP.md.

Step 3: Executing Tools & Self-Correction (RLVR Loop)

1. Retrieve relevant executable scripts or skills from 04_skills_runtime/.
2. Stream execution status and tool payloads into 05_episodic_logs/trajectories/.
3. If a step fails (reward: -1.0 or non-zero exit code):
 - * Subsection 1: Do not discard the error context.
 - * Subsection 2: Ingest the error trace from Tier 5 into active Working Memory.
 - * Subsection 3: Search 04_skills_runtime/policies/ for documented recovery heuristics.
 - * Subsection 4: Attempt remediation and log the retry attempt.

10. Complete Repository Scaffolding Script

Execute this non-interactive bash sequence to bootstrap or audit the entire 5+1 Tier cognitive filesystem topology, establishing all subdirectories, tier-level manifest.jsonl files, and root metacognitive indexes:

```
#!/usr/bin/env bash
set -euo pipefail
```

```
echo "===> Initializing 5+1 Tier Cognitive Repository Architecture..."
```

1. Scaffold all directory hierarchies

```
mkdir -p 01_raw_sources/{pdf,media,text}
mkdir -p 02_wiki_md/{concepts,architectures,entities,indexes}
mkdir -p 03_recall_cache/{jsonl,vectors,kv_store}
mkdir -p 04_skills_runtime/{prompt_skills,extracted_tools/{cli,wrappers},runtimes,policies}
mkdir -p
05_episodic_logs/{daily_driver_sync,trajectories,red_audit_sandbox,rlvr_verifiers,hygiene_reports}
```

2. Touch distributed manifest.jsonl files across all tiers

```
touch manifest.jsonl
touch 01_raw_sources/manifest.jsonl
touch 02_wiki_md/manifest.jsonl
touch 03_recall_cache/manifest.jsonl
touch 04_skills_runtime/manifest.jsonl
touch 05_episodic_logs/manifest.jsonl
```

3. Initialize Root MAP.md if absent

```
if [ ! -f "MAP.md" ]; then
```

```

cat << 'EOF' > MAP.md
# Master Repository Ontology Map

## 01. Raw Sources (Cold Archive)
- Sensory ground truth cataloged in `01_raw_sources/manifest.jsonl`.

## 02. The LLM Wiki Layer (`wiki_md/`)
- Concepts: `02_wiki_md/concepts/`
- Architectures: `02_wiki_md/architectures/`
- Entities: `02_wiki_md/entities/`
- Indexes (MOCs): `02_wiki_md/indexes/`
- Governed via OpenWiki TUI by the Files Executive Agent.

## 03. Recall Cache (High-Speed Working Memory)
- Pre-tokenized Chunks: `03_recall_cache/jsonl/`
- Vector Indices: `03_recall_cache/vectors/`
- KV Store: `03_recall_cache/kv_store/`

## 04. Procedural Runtimes, Skills & Extracted Tools
- Prompt Skills: `04_skills_runtime/prompt_skills/`
- Extracted Tools: `04_skills_runtime/extracted_tools/`
- Execution Runtimes: `04_skills_runtime/runtimes/`
- Policies & Guards: `04_skills_runtime/policies/`

## 05. Episodic Logs & Trajectories
- Daily Driver Sync: `05_episodic_logs/daily_driver_sync/`
- Trajectories: `05_episodic_logs/trajectories/`
- Red Audit Sandbox: `05_episodic_logs/red_audit_sandbox/`
- RLVR Verifiers: `05_episodic_logs/rlvr_verifiers/`
EOF
fi

echo "====> Repository structure verified and fully initialized."

```