

NovÆxorpus — the router

Tool-agnostic. Any agent/session that attaches to this repo follows this file — not just Claude Code. This is the real rules file; `CLAUDE.md` (and any other tool-specific file, e.g. a future `GEMINI.md`) is just a thin pointer to this one.

****It is not the wiki.**** No daily updates, no session notes, no running knowledge — those live in `wiki/` and `handoffs/`.

What this repo is

The pre-planning database — this project's #dumbass (Database & Universal Memory Bank Across Split Services): the universal memory layer meant to persist across every repo and every agent, not scoped to any one project. Source documents named in the 2026-08-18 session, translated verbatim into structured Markdown for the `grill-with-docs` session. Read `README.md` before doing anything.

Naming

`NAMING-CANON.md` is authoritative. NovÆxorpus · NovÆxenti · NovÆxopia · Æsop-Xi · Horizons-Ui · ÆEsc · ÆEyre. Source documents keep their original spellings — never rewrite a source to match the canon.

Hard rules are for assholes

Never stall on a tool. If it fails, do the job another way.

Layout

```
|||
|---|---|
| `01-sources/` | untouched originals + sha256 manifest |
| `02-clean/`   | verbatim cleanups |
| `03-check/`   | what the checking tool reported |
| `wiki/`       | running knowledge — built later, from `02-clean/` |
| `handoffs/`   | session state, rewritten and pushed every session |
| `unresolved.md` | durable, cross-repo, cross-agent backlog (moved here 2026-08-31) |
| `MASTER-RESUME.md`, `MASTER-CLAUDE.md` | auto-generated overview of every
attached project's RESUME.md/CLAUDE.md — see below, don't hand-edit |
| `projects/<name>/` | same-device mirror symlinks into other repos' RESUME.md/CLAUDE.md
— see below |
| `tools/`      | scripts this repo uses |
| `skills/`     | skills scoped to this repo |
```

The RESUME/unresolved/CLAUDE discipline, reflected here (2026-08-31)

Rather than this repo keeping its own separate handoff discipline, every attached project's own `RESUME.md`/`CLAUDE.md` gets mirrored here instead:

- **Same device**: `projects/<name>/RESUME.md` and `CLAUDE.md` are real symlinks into the actual repo (e.g. `projects/aesop-xi/RESUME.md` -> `~/repos/aesop-xi/RESUME.md`). Not copies — the literal same file, so there is no drift to manage.
- **A different device or cloud environment** cannot be symlinked in from here — symlinks don't cross physical machines. To attach a new environment:
 1. `git clone` this repo (novae-xorpus) into it.
 2. `git clone` whatever project repo it's working on, if not already present.
 3. Create the same local symlink within that environment.
 4. Every session there: `git pull` both repos at start, `git push` both at end (this is just hard rule 6 applied to every attached repo, not only this one).
- `unresolved.md` lives here directly (not mirrored) since it was never repo-specific to begin with.

The MASTER-*.md auto-regeneration hook

~/repos/aesop-xi/.git/hooks/post-commit` calls
`novae-xorpus/tools/regenerate\_masters.sh` after every commit in aesop-xi, which reads every `projects/\*/RESUME.md` and `projects/\*/CLAUDE.md` (via the symlinks, live working-tree content) and concatenates each into its own overview — `MASTER-RESUME.md` and `MASTER-CLAUDE.md`. This means **both reflect state as of the last commit, not the current working tree** — edit a project's file without committing, and the master stays one commit behind until the next commit fires the hook. Any other repo added under `projects/` should get the same hook installed in its own
`.git/hooks/post-commit`.

Tools and skills

tools/clean.py — the only tool here. Converts `01-sources/` to `02-clean/` in one pass: pymupdf for PDFs, pandoc for docx and html, plain read for text. Detects PDFs by magic bytes, so extensionless files work. Strips a fixed furniture list and runs a word-frequency loss check against the raw extraction. Re-run it whenever new sources land — it is idempotent and rewrites the whole tree.

python3 tools/clean.py

Its loss check does ****not**** satisfy hard rule 5. Same script, no vote.

No skills yet. Populated when the grill session defines what's needed — not before, and not by copying in a library. A skill without a stated trigger does not get installed.

Retrieval — how Drive files actually come down

The Google Drive connector works and is the normal path. Large files are not a problem: when an MCP result exceeds the inline cap the harness spools the whole thing to a JSON file on disk and returns the path. Read that file, decode the `content` field, write the bytes — the base64 never enters context. This works for multi-megabyte binaries. ****Never paste base64 through model output****; that truncates silently and has already corrupted files once.

rclone was never needed. No remote is configured and none is required.