

AGENTS.md — Operational Boundaries & Rules of Engagement: aesop-xi

Subsystem: `aesop-xi/`

Applies To: All Autonomous Agent Personas, Planning LLMs, and Execution Swarms

1. Operating Boundaries & Invariants

1. **Mandatory Clearance Tokens:**

- No agent may execute shell tools, mutate filesystem artifacts, or trigger background subdaemons without first requesting and receiving a signed `ClearanceToken` from `arbitration_engine.py`.

2. **Thermal & Resource Compliance:**

- Agents must respect `hardware_throttler.py` directives. If status is `WARM` (42°C–47.9°C), high-compute batch jobs and 9B local model passes must yield to interactive user flows. If `CRITICAL` ($\geq 48^\circ\text{C}$), only emergency recovery tasks execute.

3. **Sensory Immutability:**

- Absolute zero-modification rule on `01_raw_sources/` and raw archives across all connected repositories.

4. **No Legacy Nomenclature:**

- The term "Omni Claw" or "Omni-Claw" is permanently prohibited across all outputs, specs, and prompts. The canonical names are strictly `horizons-ui`, `novus-aexenti`, `novaecopia`, `aesop-xi`, and `novae-xorpus`.

5. **Separation of Concerns:**

- Do not mix executable tools into documentation dumps. Keep scripts in `tools/` or `arbitration/`, prompt heuristics in `skills/`, and retrieval passages in `chunks/` (`chunk.jsonl`).
-

2. Red Auditor Hand-Off Protocol

- Worker agents **MUST NOT** attempt to interact directly with `.incognito_red_sandbox/`.
- Candidate training batches or speculative script trajectories are dropped into `incoming/`.

- The background daemon `run_audit.sh` processes batches asynchronously using `sandbox_evaluator.py`.
- Output is binary: `approved/` (+1.0) or `quarantined/` (-1.0).