

Master Planning Session Synthesis: Architecture, Models & Tools

Consolidated Technical Decisions & Strategic Execution Plan

Derived from:
[CC-PLAN-SESSION-\(1-File\)/CC_Planning-session-UNORGANIZED.txt](#).
Status: Canonical Living Wiki Reference | Clean Markdown Layer
Target Repository: [data_vault/02_wiki_md/](#)

1. Executive Summary & Problem Resolution

The operator's raw planning session in [CC_Planning-session-UNORGANIZED.txt](#) outlines the core operational hurdles and answers the central architectural questions:

- Preventing Open-Weights Model Bloat:** Models do not retain weights state across sessions. Instead of letting context bloat occur through unpruned multi-turn chats, the dual-agent setup isolates tasks: the 0.8B Qwen model handles prompt restructuring, while the 9B model executes tasks using scoped mem0 habit keys and OB1 static protocols.
- GCP Cloud Optimization:** \$1,000 Vertex AI Agent Builder developer credits are channeled into heavy script compilation, database search applications, and weekly RLVR dataset fine-tuning via asymmetric IAM peering.
- The On-Device Concierge Pipeline:** Replaces manual typing with voice-to-text (Silero VAD + Moonshine) → automated meta-prompt cleaner → user approval → model execution → Kokoro TTS audio response.

2. Dual-Agent Allocation & Model Profiles

Model Role	Assigned Model	Quantization / Format	Memory Budget	Hardware Allocation	Primary Function
Triage / Concierge	Qwen 2.5 0.5B / Qwen 3.5 0.8B	GGUF Q4_0	~850 MB RAM	Snapdragon Hexagon NPU	Strips speech-to-text errors, compiles

Model Role	Assigned Model	Quantization / Format	Memory Budget	Hardware Allocation	Primary Function
					structured Markdown meta-prompts, routes queries.
Core Executor	Qwen 3.5 9B / Gemma 4 12B QAT	GGUF Q4_0 / QAT	4.8–7.5 GB RAM	Snapdragon NPU & GPU	Multi-step coding, tool execution, structural analysis, file manipulation.
Local Knowledge Hub	DeepSeek R1 Distill Qwen 8B	FP16 / GGUF Q8_0	8.0 GB RAM	NVIDIA Jetson Orin Nano Super	Background cross-checking, heavy multi-file retrieval, OB1 vector queries.
Cloud Frontier	Claude 3.7 Sonnet / Gemini 2.5 Pro	Cloud API via OmniRoute	0 MB Local RAM	GCP / OpenRouter Cloud	Complex architectural verification, code auditing, high-compute RLVR relabeling.

3. Concrete Action Schemas & Tool Ingestion Order

From Part 1 of the consolidated roadmap:

- **Priority 1:** Knowledge Curation & Universal JSONL Marker Indexing across all documentation.
- **Priority 2:** Tool Harness Ingestion (Matt Pocock Skills, Honey for Devs, Prime Agent persistent REPL, ECC, GSD, Claude Video, Code Review Graph, mem0, OmniRoute).

- **Priority 3:** Local Workspace Infrastructure & Multi-Agent Grill Sessions.
- **Priority 4:** 3-APK Native Android Suite Deployment.
- **Priority 5:** Local Dual-Model Inference & Continuous Improvement RLVR Loop.
- **Priority 6:** Multi-Node Mesh Orchestration across Node Alpha, Beta, Gamma, and Delta.