

AGENTS.md — novus-aexenti

Operational Guidelines & Model Invariants

1. Role Definitions

A. 0.8B Triage Model (Qwen3.5-0.8B-Instruct)

- **Target Hardware:** Snapdragon 8 Elite (Hexagon v79 HTP / NPU).
- **Execution Format:** GenieX / GGUF Q4_0.
- **Scope:**
 - Classifying incoming voice/text input from Horizons UI and Æyre.
 - Fast routing decisions (evaluating token counts and complexity thresholds).
 - Direct execution of atomic queries (< 120 tokens, no code generation).
- **Context Budget:** Max 4,096 tokens (output capped at 512 tokens).

B. 9B Executor Model (Qwen3.5-9B-Instruct)

- **Target Hardware:** Jetson Orin Nano Super (CUDA) or local NPU.
- **Execution Format:** llama.cpp / GGUF Q4_0.
- **Scope:**
 - Complex multi-step reasoning, tool execution, and code synthesis.
 - File-by-file refactoring and structural audits.
 - Receiving compressed meta-prompts formatted via `executor_bridge.py`.
- **Context Budget:** Max 16,384 tokens (output capped at 4,096 tokens).

2. Operational Invariants

1. **Non-Destructive Execution:** Under no circumstances should an agent overwrite or delete original source documents in cold storage (`01_raw_sources/`).
2. **Deterministic Fallbacks:** All tools and scripts invoked by agents must yield deterministic exit codes (0 for success, non-zero for error traces).
3. **Session Checkpointing:** Agents performing multi-step actions must update `active_execution_paths.json` after every step to ensure recovery on power loss or task interruption.
4. **No Hallucinated Tool Names:** Strict adherence to canonical repository names: `horizons-ui`, `novus-aexenti`, `novaecopia`, `aesop-xi`, `novae-xorpus`, and `#d.u.m.b.a.s.s..`