

AUDIT-06-NOVUS_AEXENTI_NOVAECOPIA.md

Audit & Extraction Report: Subfolder 6 of

___Lex-Novi-Æxentis-Copiæ (--• 🦁
NovusÆxenti 🌳 NovÆcopia ⚙️ • ~)

Scope: Cognitive MoE Brain (NovusÆxenti), "The Claw" Agent Swarm Harnesses (NovÆcopia), and on-device Qwen weight registries.

Parent Source: ___Lex-Novi-Æxentis-Copiæ/--• 🦁 NovusÆxenti 🌳 NovÆcopia ⚙️ • ~ (1uiLHrw-sPJAv0PgaDRI5XhCa6-pIKtWH).

Status: All original files preserved 100% untouched.

1. Purpose & System Roles: The Brain & The Claw

This folder establishes the dual-engine core:

- **NovusÆxenti (novus-aexenti/):** The Cognitive MoE Brain. Coordinates the asymmetric tandem: Qwen 0.8B intent triage vs. Qwen 9B deep reasoning, integrating with mem0 and the Reasoning Bank.
 - **NovÆcopia (novaecopia/):** "The Claw". The execution harness layer that hosts autonomous CLI swarms (OpenWiki, Hermes 3, Aider, Claude Code Android) and binds them to local MCP servers.
-

2. Major Discovery: Active On-Device Model Weights

Inside subfolder #QWEN_MODELS:

- **Qwen3.5-2B-Q4_0.gguf** (File Size: **1.21 GB**): Live, pre-quantized GGUF binary ready for direct loading on the Snapdragon 8 Elite Hexagon NPU via GenieX.
- **Mer0vin8ian/Qwen3.5-9B-GGUF**: Hugging Face repository specifications for your primary 9B executor model.
- **Qwen3-4B-Thinking-2507-IQ4_NL.gguf**: Specialized reasoning model checkpoint with explicit <think> scratchpads.

Action: Registered into MODEL_REGISTRY in npu_manager.py and codified into README_NOVUS_AEXENTI.md.

3. Subfolder Audit & Swarm Breakdown

--• 🦁 NovusÆxenti 🌳 NovÆcopia 🕸 • ~ /

|

| — 📁 1. #QWEN_MODELS/ # On-device weight binaries & quantization guides

| | — Qwen3.5-2B-Q4_0.gguf (1.21 GB) # Live binary weights for edge inference

| | — Mer0vin8ian/Qwen3.5-9B-GGUF # Primary 9B executor target

| | — Qwen3-4B-Thinking-2507-IQ4_NL # Thinking / reasoning variant

|

| — 📁 2. CLIS-AND-AGENTS/ # The Swarm Harnesses ("The Claw")

| | — open-wiki/ # LangChain OpenWiki CLI (GLM 5.2 repo cleaner)

| | — hermes/ # Nous Research Hermes 3 function-calling agent

| | — aider/ # Git-tracked autonomous code editor CLI

| | — github-mcp/ # MCP bridge for GitHub repository operations

| | — claude-code-android/ # Termux/Android integration for Claude Code

| | — Linxr and pockr/ # Matt Pocock developer skills framework

|

| — 📁 3. On-device_Model-Guide/ # RAM boundaries (keeping active weights under 5GB)

| — 📁 4. Agent-Builds&Guidelines/ # Autonomous persona directives & tool schemas

| — 📁 5. #GOOGLE _AGENTIC-AI/ # Vertex AI Agent Builder cloud fallback configs

4. What Was Combined & What Was Trimmed

- **Fluff Trimmed:** Bypassed duplicate Google Search PDF dumps ([Qwen3-4B-Thinking... - Google Search.pdf](#) [6.2 MB]). The underlying quantization parameters were extracted into [MODEL_REGISTRY](#).
- **Combined Harnesses:**
 - Synthesized Hermes 3, Aider, and OpenWiki into the unified modular harness directory under [novaexopia/modular_harnesses/](#).
 - Unified the Qwen 2B and 9B configurations into the dual-agent routing logic in [01_SOVEREIGN_NODE_AND_APK_TOPOLOGY.md](#).