

02_wiki_md/	# 5+1 outer tier stays; INTERIOR is properly extensive
— MAP.md	# Read this FIRST (Part 1.3 directive)
— manifest.jsonl	# Queried by RAG; never loaded whole
— vendors/	# BY-VENDOR — first-class, each vendor gets a full
sub-tree	
— MAP.md	
— manifest.jsonl	
— google/	# ★ MAJOR — dedicated cut per operator directive
— MAP.md	
— platform/	# GCP: Vertex AI (formerly), GenAI Agent Platform, Cloud
Shell,	
	Developer Suite, Cloud Dashboard, Anti-Gravity, Gemini Spark
— android/	# NDK: ASharedMemory, SCM_RIGHTS, AndroidManifest
process isolation,	
	Foreground Service / LMK, Media SDK, Video Game SDK,
Accessibility	
— models/	# Gemma models — Gemma 4 12B (Node Beta), Gemma
4 E4B, PiperTTS-DE	
— gemini/	# Gemini Nano (Chrome), Gemini Spark, Gemini API, AI
Studio	
— notebooklm/	# NotebookLM + notebooklm-py
— jetpack-compose/	# Native Android UI framework (candidate for
Horizons UI rebuild)	
— credits-and-accounts/	# \$1000 Vertex credits, Business account, AI Pro
extended tier,	
	Developer account, beta/demo access log
— runbooks/	# GCP cross-account handshake, IAM setup,
service-account provisioning	
— qualcomm/	# ★ MAJOR
— MAP.md	
— qairt-sdk/	# QAIRT SDK docs, headers, HTP specs
— genie-bundles/	# Pre-compiled Genie bundles from QAI Hub
— hexagon/	# NPU v79/v73, INT8 TOPS, FastRPC, DMA-BUF
— qat/	# Quantization Aware Training — operator uses in
conjunction w/ GenieX	
— models/	# Qualcomm-hosted models (Qwen3.5-0.8B QAI,
LiteHRNet, Phi-4-Mini)	
— nvidia/	# ★ MAJOR
— MAP.md	
— jetson/	# Orin Nano Super, CUDA cores, 60-70 TOPS
— cuda/	# CUDA runtime + toolchain

models/	# Parakeet (STT), NVIDIA Polar, etc
github/	# ★ MAJOR — operator has forks + custom repos
— MAP.md	
— owned/	# c10vis-poem repos (novae-xorpus, aesop-xi, openwiki)
fork,	
—	claude-code-android, OB1, ECC-aesop)
— forked/	# Everything he's forked from upstream
— watched/	# Upstream repos he tracks (Anthropic MCP, Prime)
Agent, Graphify,	
—	obsidian-skills, mattpocock/skills, PrimeIntellect/prime-agent)
— actions/	# GitHub Actions workflows, CI/CD configs
— anthropic/	
— MAP.md	
— claude-code/	# Claude Code CLI docs
— mcp/	# Anthropic MCP registry, official server list
— plugin-marketplace/	# The now-agent-friendly plugin repo
— api/	# Messages API, tool use, prompt caching
— primeintellect/	
— MAP.md	
— prime-agent/	# RLM harness, Continual Harness, /refine
— env-hub/	
— deepseek/	# NEW — coming harness per operator
— MAP.md	
— harness/	# DeepSeek harness spec (to be filled in)
— models/	
— unsloth/	# Fine-tuning docs, GGUF export, quantization guides
— liquid-ai/	# LFM2.5-8B-A1B reference
— together/	# OSCAR (2-bit attention-aware KV cache) research
— mattpocock/	# mattpocock/skills (grill-with-docs, etc)
— kepano/	# obsidian-skills
— affaan-m/	# ECC (Everything Claude Code)
— graphify-labs/	# Graphify (AST → knowledge graph)
— openrouter/	# OpenRouter API, GLM 5.2 access, model catalog
— z-ai/	# GLM 5.2 model docs, pricing, 1M context notes
— weights/	# BY-MODEL — every model weight file catalogued
— MAP.md	
— manifest.jsonl	# per-weight: repo id, quantization, size, on-disk-path,
target-hardware	

- active/ # Weights currently loaded/loadable on-device
 - qwen-3.5-0.8b-qai-hub-genie/ # Executor (NPU-pinned via QAIRT)
 - qwen-3.5-9b-unsloth-gguf-q4_0/ # Query (NPU via GenieX)
 - qwen-3.5-2b-q4_0-gguf/ # ← ALREADY ON DISK per Lex-Novi discovery
- (1.21 GB)
- candidates/ # Weights to evaluate but not yet chosen
 - gemma-4-12b-it-qat-ud-q4_k_xl/ # 12B for Node Beta (Jetson, 8GB tight) — per operator plan
- gemma-4-e4b-qai-hub/ # E4B (Qualcomm-hosted)
 - phi-4-mini-qai-hub/ # Help-desk agent tier
 - granite-mini/ # Help-desk agent tier
 - qwen-3.5-4b/ # Fallback tier if 9B chokes
 - twin-9b/ # (verify spelling — may be a mishear for a different 9B)
 - rejected/ # Weights evaluated and dropped, with reason
 - pocket-35b-gguf/ # >8GB, impossible on-device (operator decision)
- 2026-09-05)
- voice/ # Voice-specific model weights
 - moonshine-small-onnx/ # STT
 - kokoro-82m/ # TTS
 - sherpa-onnx/ # TTS alt
 - silero-vad/ # VAD (technically an ONNX classifier, but voice-stack)
 - parakeet/ # NVIDIA STT alt
 - vision/ # Vision model weights
 - litehrnet/ # Qualcomm-hosted vision
- runtimes/ # BY-RUNTIME — the low-level executors
 - MAP.md
 - qairt/ # libQnnHtp.so + env vars + genie-t2t-run CLI
 - genie-x/ # Wraps llama.cpp + GGML Hexagon backend + CLI
 - llama-cpp/ # Base llama.cpp + GGML Hexagon backend +
- haozixu/llama.cpp-npu fork
 - onnx-runtime/ # ORT + QNN Execution Provider for NPU
 - whisper-cpp/ # Voice STT runtime
 - cuda/ # CUDA on Jetson (Node Beta)
 - node-js/ # Node.js LTS runtime for MCP servers + OpenWiki + web
- tooling
- engines/ # BY-ENGINE — higher-level wrappers around runtimes
 - MAP.md
 - genie-x-engine/ # (same as runtimes/genie-x but with the
- scheduler/serving surface)
 - moonshine-onnx-engine/ # STT engine
 - kokoro-tts-engine/ # TTS engine
 - piper-tts-engine/ # Alt TTS engine

		silero-vad-engine/	# VAD engine
		llama-server/	# llama.cpp's HTTP server wrapper
		openwiki-tui-engine/	# OpenWiki TUI (multi-profile: file_administrator,
	oeracle_helpdesk)		
		harnesses/	# BY-HARNESS — outer control loops
		MAP.md	
		ecc/	# Everything Claude Code — governs Claude Code CLI
	sessions		
		overview.md	
		internal-routing.md	# honey-crush / nexus-mapper / ecc-planner /
	px-reader		
		plugin-marketplace.md	
		skill-catalog.md	# 286 skills, 68 agents
		prime-agent/	# Governs on-device open-weight models via RLM
		overview.md	
		rlm-loop.md	# Recursive Language Model + persistent Python REPL
		continual-harness.md	# State refinement + rollback
		openai-client-config.md	# Points to OmniRoute localhost:20128
		deepseek/	# ★ NEW — coming per operator
		overview.md	# (TBD)
		coexistence-with-prime.md	
		gsd-v2/	# Build-phase execution engine (fresh context per task, git
	branch-per-slice)		
		pocock-skills/	# Planning layer (grill-with-docs, to-spec, to-tickets, TDD,
	code-review)		
		honey-for-devs/	# Dev-focused skill harness
		aider/	# Aider workspace agent
		smol-agents/	# Low-overhead script-fragment generator pods
		nano-agents/	# (per operator Google docs on nano agent building)
		turbo-quant-swarms/	# Quantized reasoning models for high-speed traffic
		agents/	# BY-AGENT — specific personas running under harnesses
		MAP.md	
		manifest.jsonl	# per-agent: harness, model, memory-scope,
	tool-allowance, on-node		
		oeracle/	# Oeracle — on-device help-desk oracle (Qwen 9B on
	Hexagon)		
		file-administrator/	# Novus/Exenti-Housekeeper (GLM-5.2 file admin)
		cross-auditor/	# Home-node public auditor (script collector + log
	aggregator)		
		red-auditor/	# ⚠ INCOGNITO — link only; actual spec lives in ~/.red/
	per §6 isolation		
		early-trend-scraper/	# Daily scraper against dev news + GitHub trends

- | |— home-assistant-auditor/ # Home Assistant cross-agent auditor
- | |— it-helpdesk/ # IT Help Desk & Manual Operator
- | |— web-ingestion-monitor/ # Tracks upstream model/tool/repo changes
- | |— npu-inference-manager/ # Balances weight offloading Alpha↔Beta↔Cloud
- | |— executor-0.8b/ # 0.8B executor persona (fast task exec)
- | |— query-9b/ # 9B query persona (meta-prompt compilation, cloud tool
- calls)
 - | |— frontier-callers/ # Personas allowed to call frontier APIs (permission gated)
 - | |— claude-caller.md
 - | |— glm-caller.md
 - | |— gemini-caller.md
 - | |— fable-caller.md
 - | |— opus-caller.md
- |— protocols/ # BY-PROTOCOL — specs (not code)
 - | |— MAP.md
 - | |— aesop-xi/ # THE ethical arbitration protocol (Agentic Execution Split
- Ops)
 - | |— mcp/ # Model Context Protocol spec
 - | |— manifest-jsonl-schema/ # Universal JSONL marker schema (Part 1 spec)
 - | |— skill-onboarding-schema/ # AESOP_XI_Skill_Onboarding_Template JSON
- schema
 - | |— w5h-matrix/ # Who/What/When/Where/Why/How operational lens
 - | |— rlvr/ # Reinforcement Learning with Verifiable Rewards
 - | |— kag-recurse/ # Knowledge-Augmented-Generation recursive loop spec
 - | |— rag-pipeline/ # Retrieval-Augmented-Generation pipeline (see Part 6)
 - | |— adb-loopback/ # The "laptop trick" — ADB-to-WebSocket / UNIX-host
- tunnel
 - | |— unix-socket-ipc/ # AF_UNIX + SCM_RIGHTS FD-passing between
- daemons
 - | |— extraction/ # DATA EXTRACTION protocols (see Part 3.1 definition
- below)
 - | |— omniroute-extraction.md # OmniRoute AS the memory extraction layer
 - | |— skill-tool-dual-extraction.md # From docs → Skills + Tools
 - | |— voice-extraction.md # STT extraction pipeline
 - | |— memory-pipeline-order/ # OmniRoute → Reasoning Bank → OB1 → mem0
- daemons/models
 - | |— droiddesk-workstation/
 - | |— incognito-red-audit/ # 7-invariant isolation contract (from 01B §6)
 - | |— cost-checking/ # Cost audit protocol (see Part 7)
- |— scripts/ # BY-EXECUTABLE — actual scripts (mirror of
- 04_skills_runtime/extracted_tools/)
 - | |— MAP.md

- setup-aesop-sh.md # doc explaining the setup script
- convert-raw-to-markdown-py.md
- generate-jsonl-markers-py.md
- gcp-cross-account-handshake-sh.md
- run-audit-sh.md # Red Auditor daemon (documented here, code lives in
- ~/.red/)
- boot-sh.md # Reboot execution flow (boot.sh → tmux llm → ow)
- system-housekeeper-sh.md
- agent-panel-sh.md
- npu-manager-py.md
- compile-manifest-py.md
- doc-to-skill-and-tool-py.md
- export-drive-docs-js.md
- log-builder-py.md
- tools/ # BY-CAPABILITY — reusable tool descriptions (5-doc-type: Tool)
- MAP.md
- (one .md per tool available to any agent — mirrors extracted_tools/)
- skills/ # BY-SKILL — reusable skill descriptions (5-doc-type: Skill)
- MAP.md
- (one .md per SKILL.md — mirrors prompt_skills/)
- projects/ # BY-PROJECT — per-repo subvaults
- MAP.md
- aesop-xi/
- novus-aexenti/
- novaexopia/
- horizons-ui/ # OR: horizons-ui-legacy/ + horizons-ui-modern/ (salvage
- split)
- aesc/ # Terminal daemon APK
- aeyre/ # Media daemon APK
- novaexorpus/ # Self-referential — the vault about the vault
- openwiki-fork/ # c10vis-poem/openwiki
- OB1/
- claude-code-android/
- ECC-aesop/
- entities/ # BY-STRUCTURAL-ENTITY — registries, schemas,
- hardware profiles
- MAP.md
- nodes/ # Alpha (Moto Razr) / Beta (Jetson) / Gamma (Rubik Pi) /
- Delta (GCP)

— schemas/	# JSONL, SKILL.md, YAML profile, JSON-RPC audit
schemas	
— canon/	# NAMING-CANON, brand ligature, motto
— permissions/	# OS permissions, ADB rights, tool guards, security
clearance	
— architectures/	# BY-BLUEPRINT — system diagrams, data flows,
integration maps	
— MAP.md	
— 3-apk-native-topology/	# Horizons + ÆEsc + ÆEyre with WebSocket + UNIX
socket bridges	
— 4-node-hardware-mesh/	# Alpha/Beta/Gamma/Delta
— memory-pipeline/	# OmniRoute → Reasoning Bank → OB1 → mem0
(with KAG + RAG loops)	
— rlvr-flywheel/	# Trajectory → Red Audit → GCP → weights back to edge
— laptop-trick/	# ADB loopback / debug-mode escape
— droiddesk-workstation/	
— modular-hot-swap/	# ★ NEW — spec for swapping agent harness without
cascading changes (Part 8)	
— memory-subsystem/	# BY-MEMORY-COMPONENT — first-class, gets its
own top-level cut	
— MAP.md	
— omniroute/	# ★ AS MEMORY EXTRACTION LAYER
(localhost:20128) — see Part 5 correction	
— reasoning-bank/	# Multi-model execution ledger + crash recovery
— ob1/	# Postgres vector protocol (Open Brain), ground-truth
retrieval	
— mem0/	# In-session episodic state + user habits
— sqlite/	# Embedded on-device KV + relational (sub-5ms)
— postgres/	# Authoritative relational + vector on Node Beta
— kv-cache/	# ★ NEW — KV cache spec vs SQL (see Part 6)
— continual-harness/	# Reset-free online adaptation
— graphify/	# AST + knowledge graph mapping
— code-review-graph/	# Cross-agent graph code checking
— notebooklm-py/	# Deep analytical query engine
— honey-for-devs-mem/	# Any memory-hooks Honey exposes
— nope-databank/	# Red Auditor's negative reference set (link only; lives in
~/.red/)	
— runbooks/	# BY-OPERATIONAL-PROCEDURE — step-by-step
how-tos	
— MAP.md	
— glm-openwiki-setup.md	

— droiddesk-install.md	# ← Tonight's first target
— aesc-daemon-bootstrap.md	# ← Tonight's second target
— horizons-salvage-plan.md	# ★ NEW — salvage philosophy applied (Part 10)
— gcp-cross-account-handshake.md	
— qairt-npu-proof-of-life.md	
— geniex-9b-npu-test.md	
— car-wash-integration-test.md	
— reboot-boot-sequence.md	# boot.sh → tmux llm → ow flow
— references/	# BY-EXTERNAL-SOURCE — vendor docs + papers as
Tier-2 notes	
— MAP.md	
— continual-harness-paper/	# arXiv 2605.09998v1
— together-oscar/	# 2-bit attention-aware KV cache
— nvidia-polar/	# GRPO training
— stanford-ai-index/	
— liquid-ai-lfm25/	
— memo-memory-model/	
— hexo-labs-sia/	
— omnivoice-studio/	
— futo-keyboard/	
— llm-wiki-pattern/	# Karpathy pattern
— operator-log/	# BY-SPEAKER-ATTRIBUTION (hard rule 4)
— MAP.md	
— user-corrections-log.md	# S4-style verbatim annotations
— naming-canon-history.md	
— priority-reorderings.md	
— decisions-log.md	# Every hard/soft decision with date + rationale
— unresolved-questions.md	# Cross-vault open items
— indexes/	# BY-CROSS-CUTTING-VIEW — MOCs re-slicing
everything	
— MAP.md	
— by-doc-type.md	# 5 doc types: Tool / Skill / Reference / Memory / Data
— by-vendor.md	# rollup of vendors/
—	